



# INFO20003 Database Systems

Assignment Project Exam Help

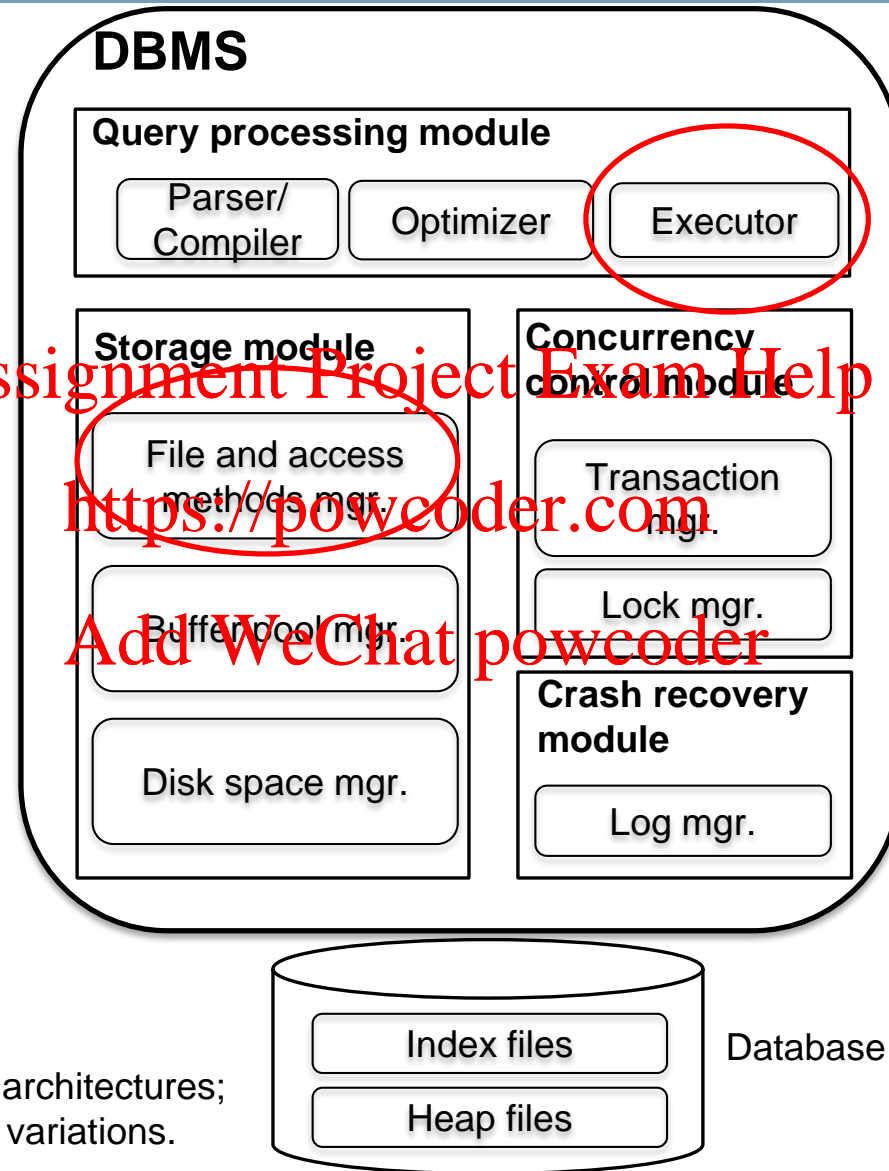
<https://powcoder.com>

Dr. Renata Borovica-Gajic  
Add WeChat powcoder

Lecture 11  
Query Processing Part I

# Remember this? Components of a DBMS

MELBOURNE



**TODAY &  
Next time**

**Will briefly  
touch upon ...**

Assignment Project Exam Help

<https://powecoder.com>

Add WeChat powecoder

This is one of several possible architectures; each system has its own slight variations.

Database



MELBOURNE

- Query Processing Overview
- Selections
- Projections

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 12 and 14, Ramakrishnan & Gehrke, Database Systems*



MELBOURNE

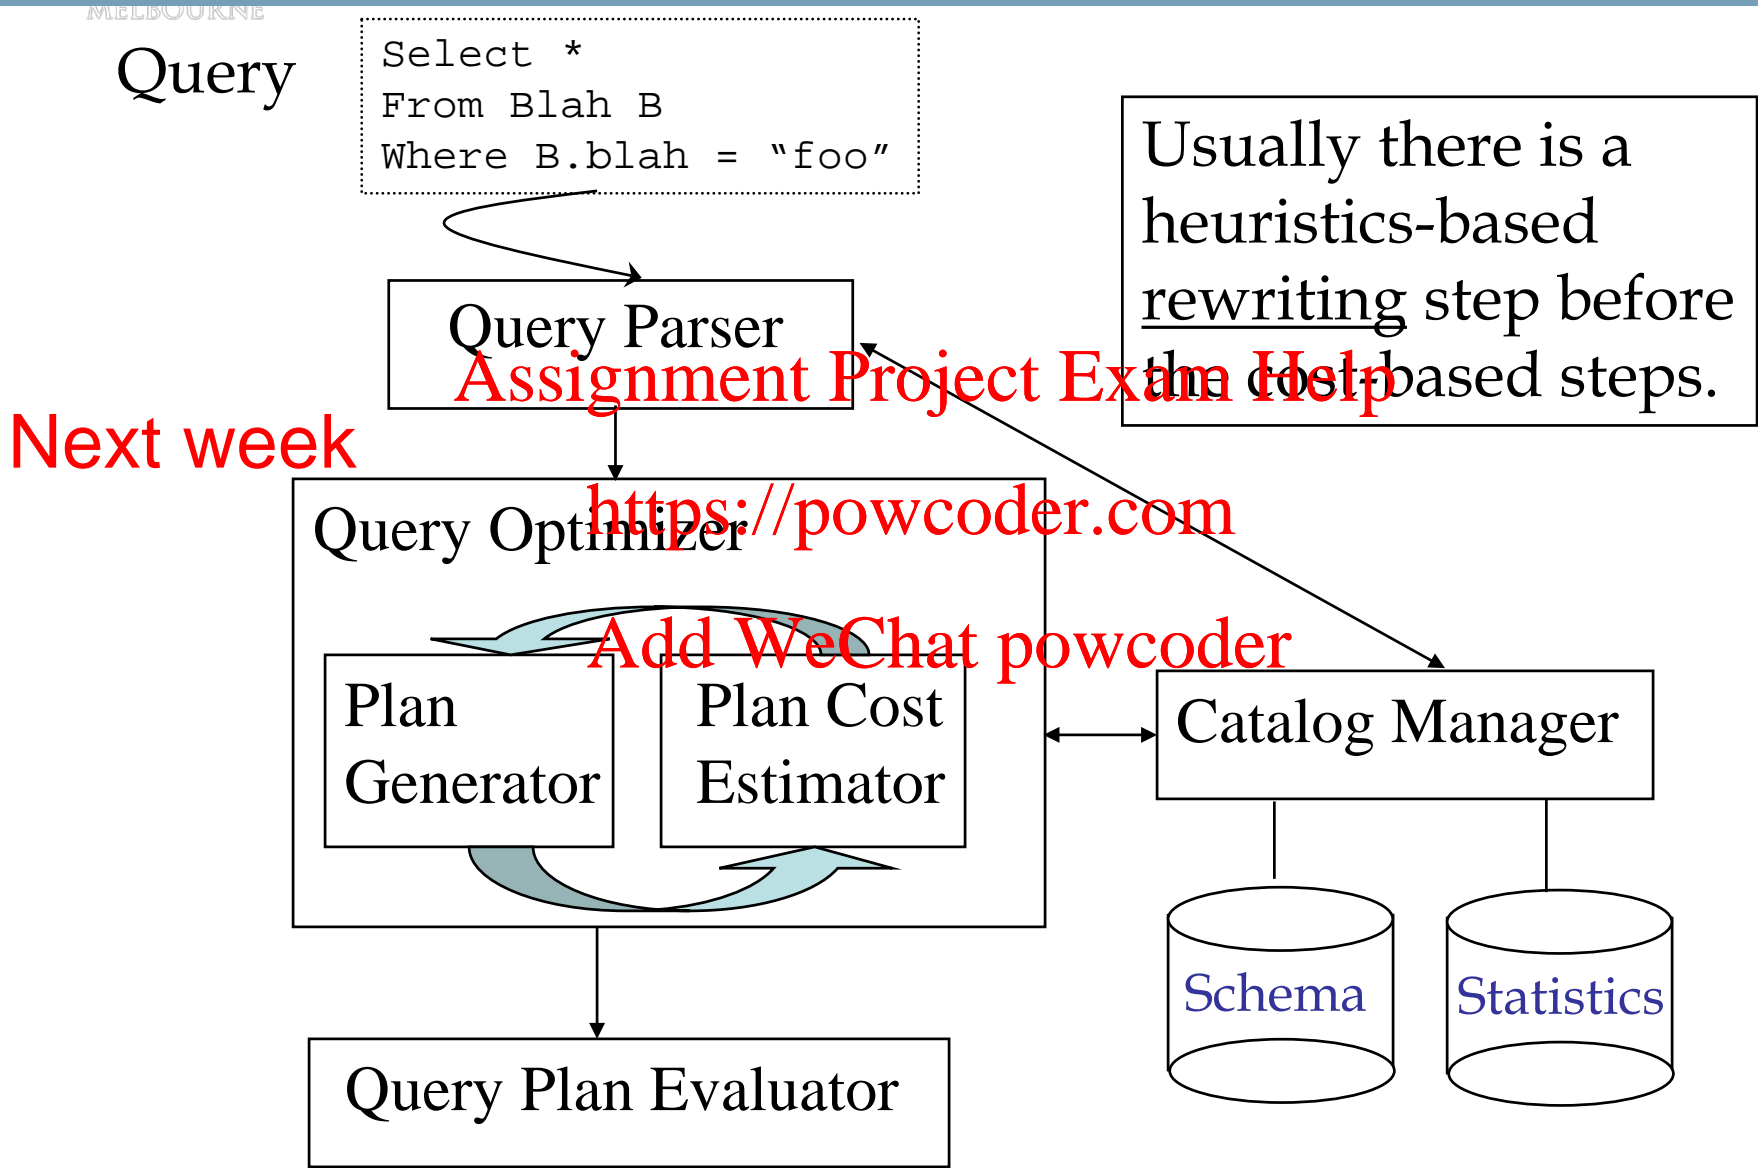
- Some database operations are **EXPENSIVE**
- DBMSs can greatly improve performance by being ‘smart’
  - e.g., can speed up 1,000,000x over naïve approach

## Assignment Project Exam Help

- Main weapons are:
  1. clever implementation techniques for operators
  2. exploiting ‘equivalencies’ of relational operators
  3. using cost models to choose among alternatives

<https://powcoder.com>

Add WeChat powcoder



- We will consider how to implement:
  - Selection ( $\sigma$ ) Selects a subset of rows from relation
  - Projection ( $\pi$ ) Deletes unwanted columns from relation
  - Join ( $\bowtie$ ) Allows us to combine two relations
- Operators can be then be <https://powcoder.com> creating *query plans*  
Add WeChat powcoder



MELBOURNE

- Query Processing Overview
- Selections
- Projections

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

Sailors (sid: integer, *sname*: string, *rating*: integer, *age*: real)  
Reserves (sid: integer, bid: integer, day: dates, *rname*: string)

- **Sailors (S):** Assignment Project Exam Help

- Each tuple is 50 bytes long, 80 tuples per page, **500 pages**

- $N = NPages(S) = 500$ ,  $p_S = NTuplesPerPage(S) = 80$

- $NTuples(S) = 500 * 80 = 40000$

<https://powcoder.com>  
Add WeChat powcoder

- **Reserves (R):**

- Each tuple is 40 bytes long, 100 tuples per page, **1000 pages**

- $M = NPages(R) = 1000$ ,  $p_R = NTuplesPerPage(R) = 100$

- $NTuples(R) = 100000$





- Of the form  $\sigma_{R.attr \text{ op } value} (R)$

- Example:

```
SELECT *  
FROM Reserves R  
WHERE R.BID > 20;
```

Assignment Project Exam Help

- The best way to perform a selection depends on:
  1. available indexes/access paths
  2. expected **size of the result** (number of tuples and/or number of pages)

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

- **Size of result** approximated as:

*size of relation \*  $\prod$  (reduction factors)*

- **Reduction factor** is usually called **selectivity**. It estimates what portion of the relation will qualify for the given predicate, i.e. satisfy the given condition.
  - This is estimated by the optimizer (will be taught next week)
  - E.g. 30% of records qualify, or 5% of records qualify

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



## 1. With no index, unsorted:

- Must scan the whole relation, i.e. perform Heap Scan
- **Cost = Number of Pages of Relation, i.e.  $NPages(R)$**
- **Example:** Reserves cost( $R$ ) = 1000 IO (1000 pages)

[Assignment Project Exam Help](#)

## 2. With no index, but file is sorted:

- cost = **binary search cost + number of pages containing results**
- **Cost =  $\log_2(NPages(R)) + (BF * NPages(R))$**
- **Example:** Reserves cost( $R$ ) = 10 I/O + ( $RF * NPages(R)$ )

## 3. With an index on selection attribute:

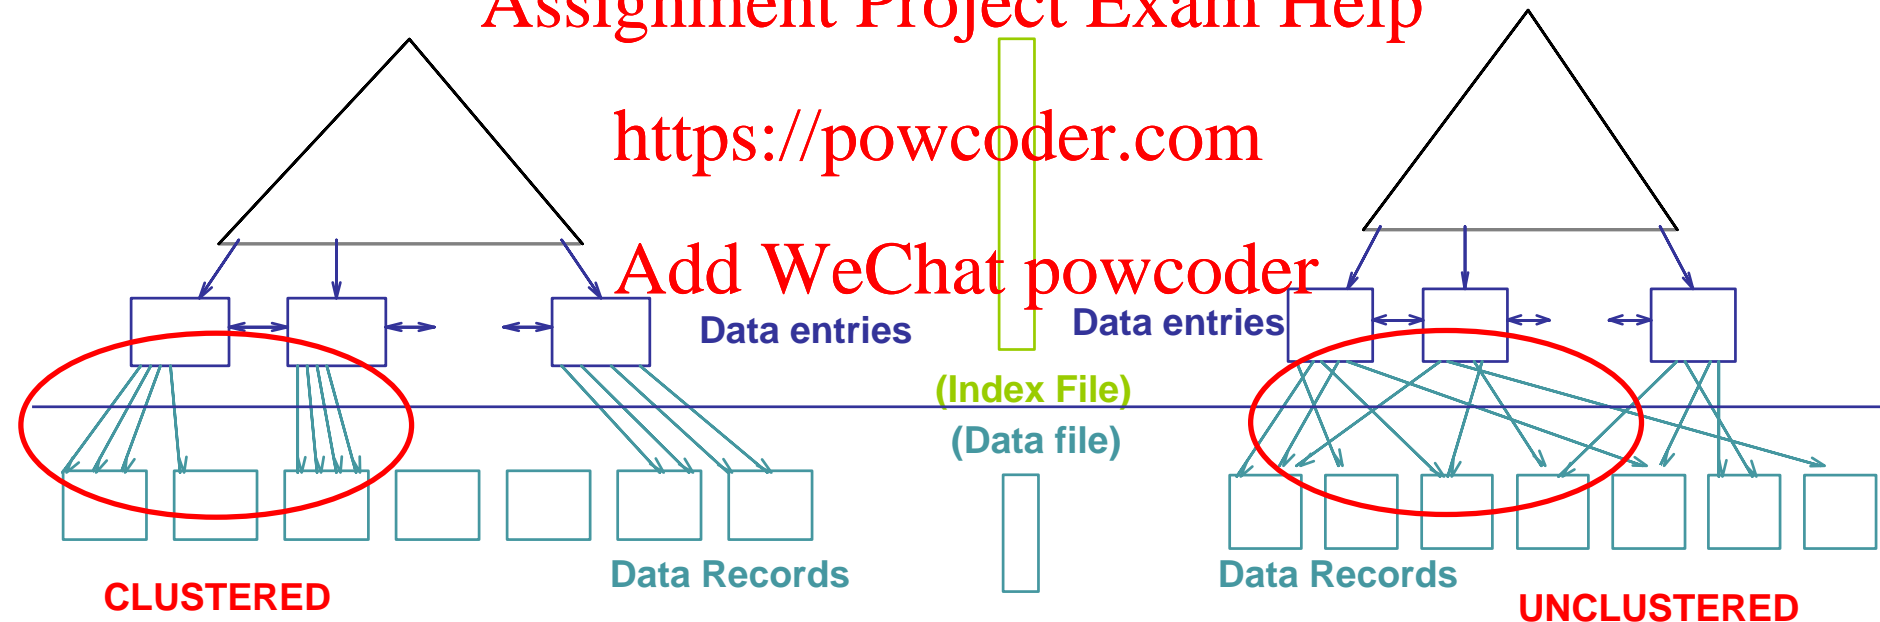
- Use index to find qualifying data entries,
- Then retrieve corresponding data records
- Discussed next....

## Clustered vs. unclustered

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

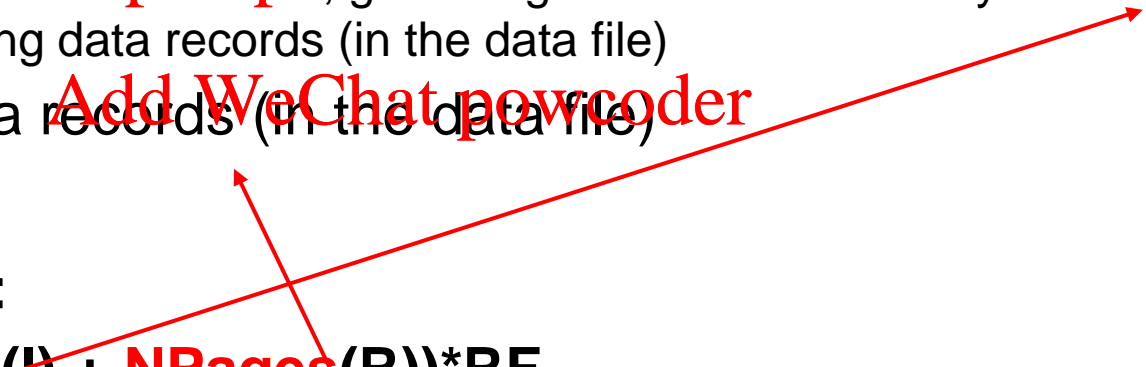


- Cost depends on the number of qualifying tuples
- Clustering is important when calculating the total cost
- Steps to perform:
  1. Find qualifying data entries:
    - Go through the index: height typically small, 2-4 I/O in case of B+tree, 1.2 I/O in case of hash index (*negligible* if many records retrieved)
    - Once data entries are reached, go through data entries one by one and look up corresponding data records (in the data file)
  2. Retrieve data records (in the data file)
- **Cost:**
  1. Clustered index:  
**Cost = (NPages(I) + NPages(R))\*RF**
  2. Unclustered index:  
**Cost = (NPages(I) + NTuples(R))\*RF**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Cost = (NPages(I) + NPages(R))\*RF

Cost = (NPages(I) + NTuples(R))\*RF

Cost = (NPages(I) + NTuples(R))\*RF

MELBOURNE

- **Example:** Let's say that 10% of Reserves tuples qualify, and let's say that index occupies 50 pages
- $RF = 10\% = 0.1$ ,  $NPages(I) = 50$ ,  $NPages(R) = 1000$ ,  $NTuplesPerPage(R) = 100$

## Assignment Project Exam Help

- **Cost:**

<https://powcoder.com>

1. Clustered index:

$$\text{Cost} = (NPages(I) + NPages(R)) * RF$$

$$\text{Cost} = (50 + 1000) * 0.1 = 105 \text{ (I/O)}$$

Cheapest access path

2. Unclustered index:

$$\text{Cost} = (NPages(I) + NTuples(R)) * RF$$

$$\text{Cost} = (50 + 100000) * 0.1 = 10005 \text{ (I/O)}$$

3. Heap Scan:

$$\text{Cost} = NPages(R) = 1000 \text{ (I/O)}$$

MELBOURNE

- Typically queries have multiple predicates (conditions)
- **Example:**  $\text{day} < 8/9/94 \text{ AND } \text{rname} = \text{'Paul'} \text{ AND } \text{bid} = 5 \text{ AND } \text{sid} = 3$
- A B-tree index **matches** (a combination of) predicates that involve only attributes in a **prefix of the search key**
  - Index on  $\langle a, b, c \rangle$  matches predicates on:  $(a, b, c)$ ,  $(a, b)$  and  $(a)$
  - Index on  $\langle a, b, c \rangle$  matches  $a = 5 \text{ AND } b = 3$ , but will not be used to answer  $b = 3$
  - This implies that only reduction factors of predicates that are **part of the prefix** will be used to determine the cost (they are called matching predicates (or primary conjuncts))

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



1. Find the **cheapest access path**
  - An index or file scan with the **least estimated page I/O**
2. Retrieve tuples using it
  - **Predicates that match** this index reduce the number of tuples *retrieved (and impact the cost)*
3. Apply the predicates that **don't match** the index (if any) later on
  - These predicates are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched (nor the total cost)
  - In this case selection over other predicates is said to be done “on-the-fly”

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





# Cheapest Access Path: Example

- **Example:**  $\text{day} < 8/9/94$  AND  $\text{bid}=5$  AND  $\text{sid}=3$
- A **B+ tree** index on **day** can be used;
  - $\text{RF} = \text{RF}(\text{day})$
  - Then,  $\text{bid}=5$  and  $\text{sid}=3$  must be checked for each retrieved tuple *on the fly*
- Similarly, a **hash index** on **<bid, sid>** could be used;
  - $\text{RF} = \text{RF}(\text{bid}) * \text{RF}(\text{sid})$
  - Then,  $\text{day} < 8/9/94$  must be checked *on the fly*
- How about a B+tree on <lname,day>? (Y/N)
- How about a B+tree on <day, lname>? (Y/N)
- How about a Hash index on <day, lname>? (Y/N)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

- Overview
- Selections
- Projections

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*



MELBOURNE

- Issue with projection is removing **duplicates**

```
SELECT DISTINCT R.sid, R.bid  
FROM Reserves R
```

**Assignment Project Exam Help**

- Projection can be done based on **hashing** or **sorting**

**Add WeChat powcoder**

**<https://powcoder.com>**

- Basic approach is to use **sorting**
  - 1. Scan R, extract only the **needed** attributes
  - 2. Sort the result set (typically using external merge sort)
  - 3. Remove **adjacent** duplicates

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

11,80
12,10
12,10
12,75
13,20
13,20
13,75

# External Merge Sort


- If data does not fit in memory do several passes
- Sort runs: Make each B pages sorted (called runs)
- Merge runs: Make multiple passes to merge runs
  - Pass 2: Produce runs of length  $B(B-1)$  pages
  - Pass 3: Produce runs of length  $B(B-1)^2$  pages
  - ...
  - Pass P: Produce runs of length  $B(B-1)^P$  pages

We will let you know

how many passes there are

Assignment Project Exam Help

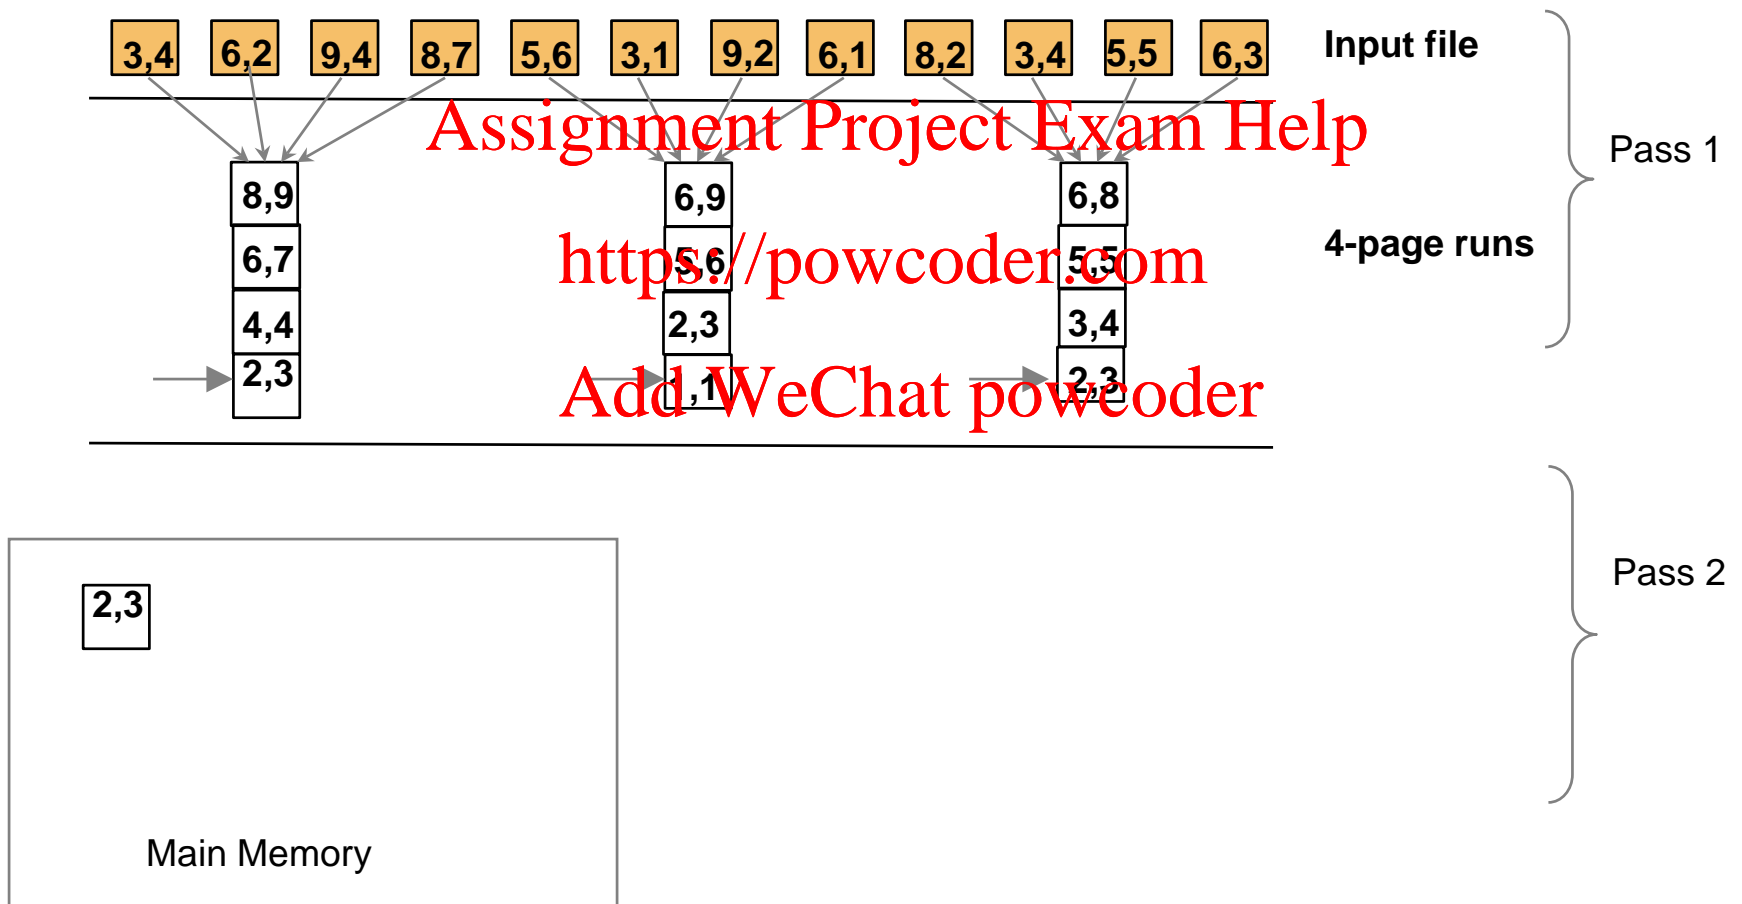
<https://powcoder.com>



Readings: Chapter 13, Ramakrishnan & Gehrke, Database Systems

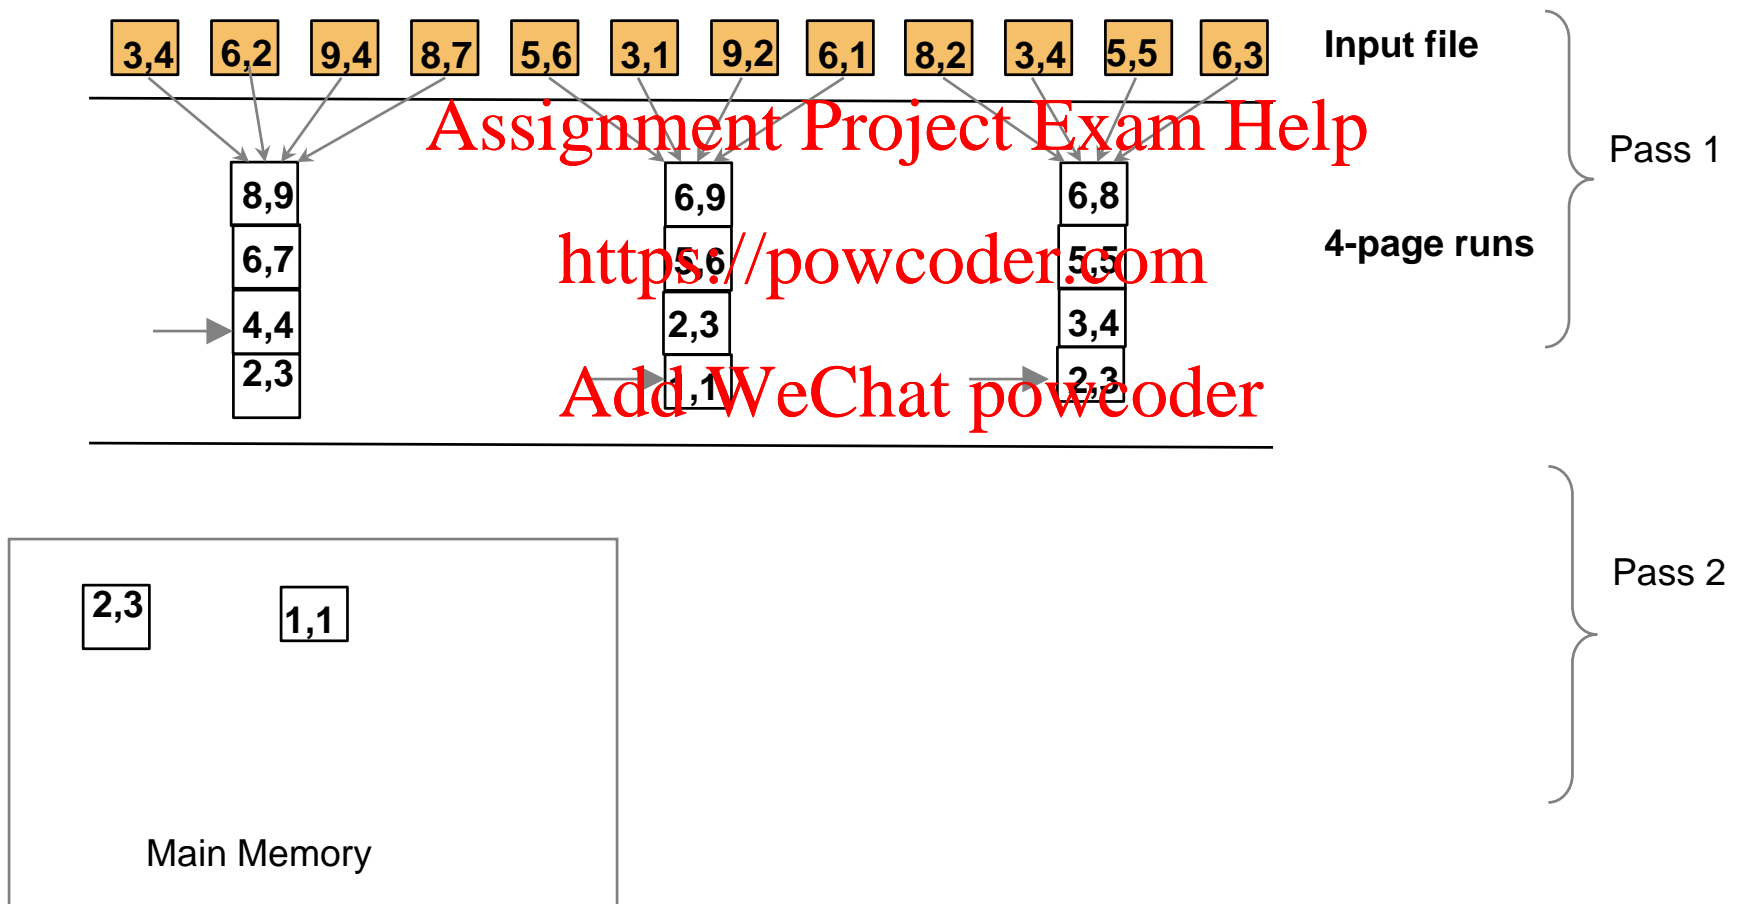
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records,  
sorting on a single attribute (just showing the attribute value)




# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



# External Merge Sort: Example

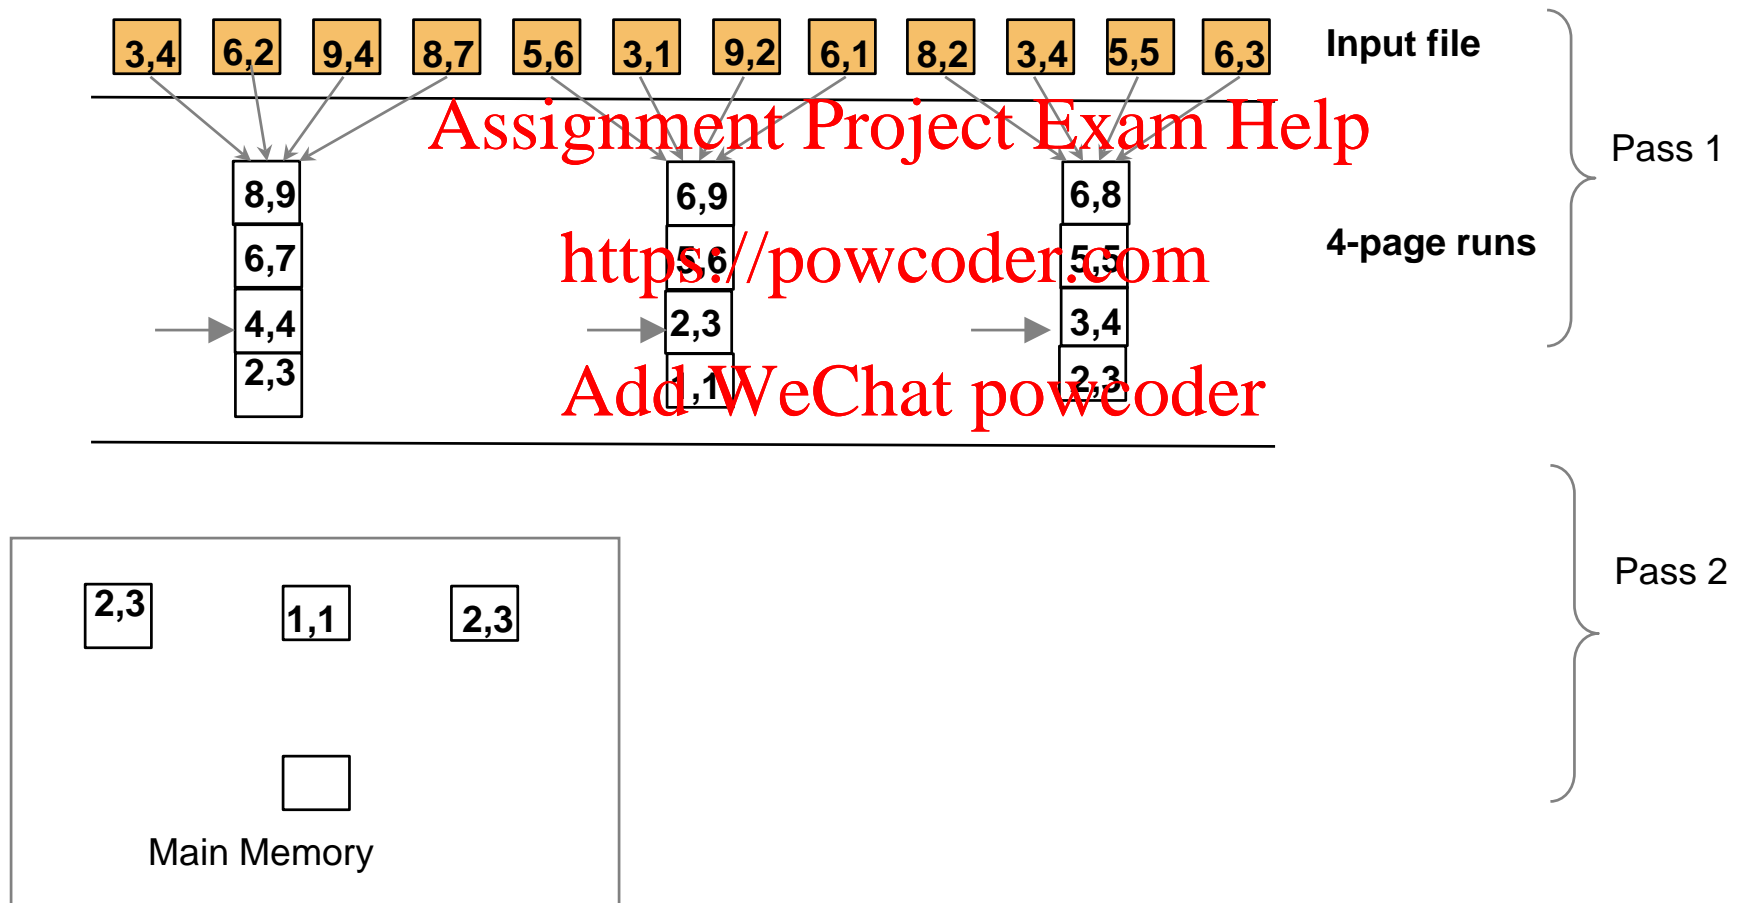
# buffer pages in memory  $B = 4$ , each page 2 records





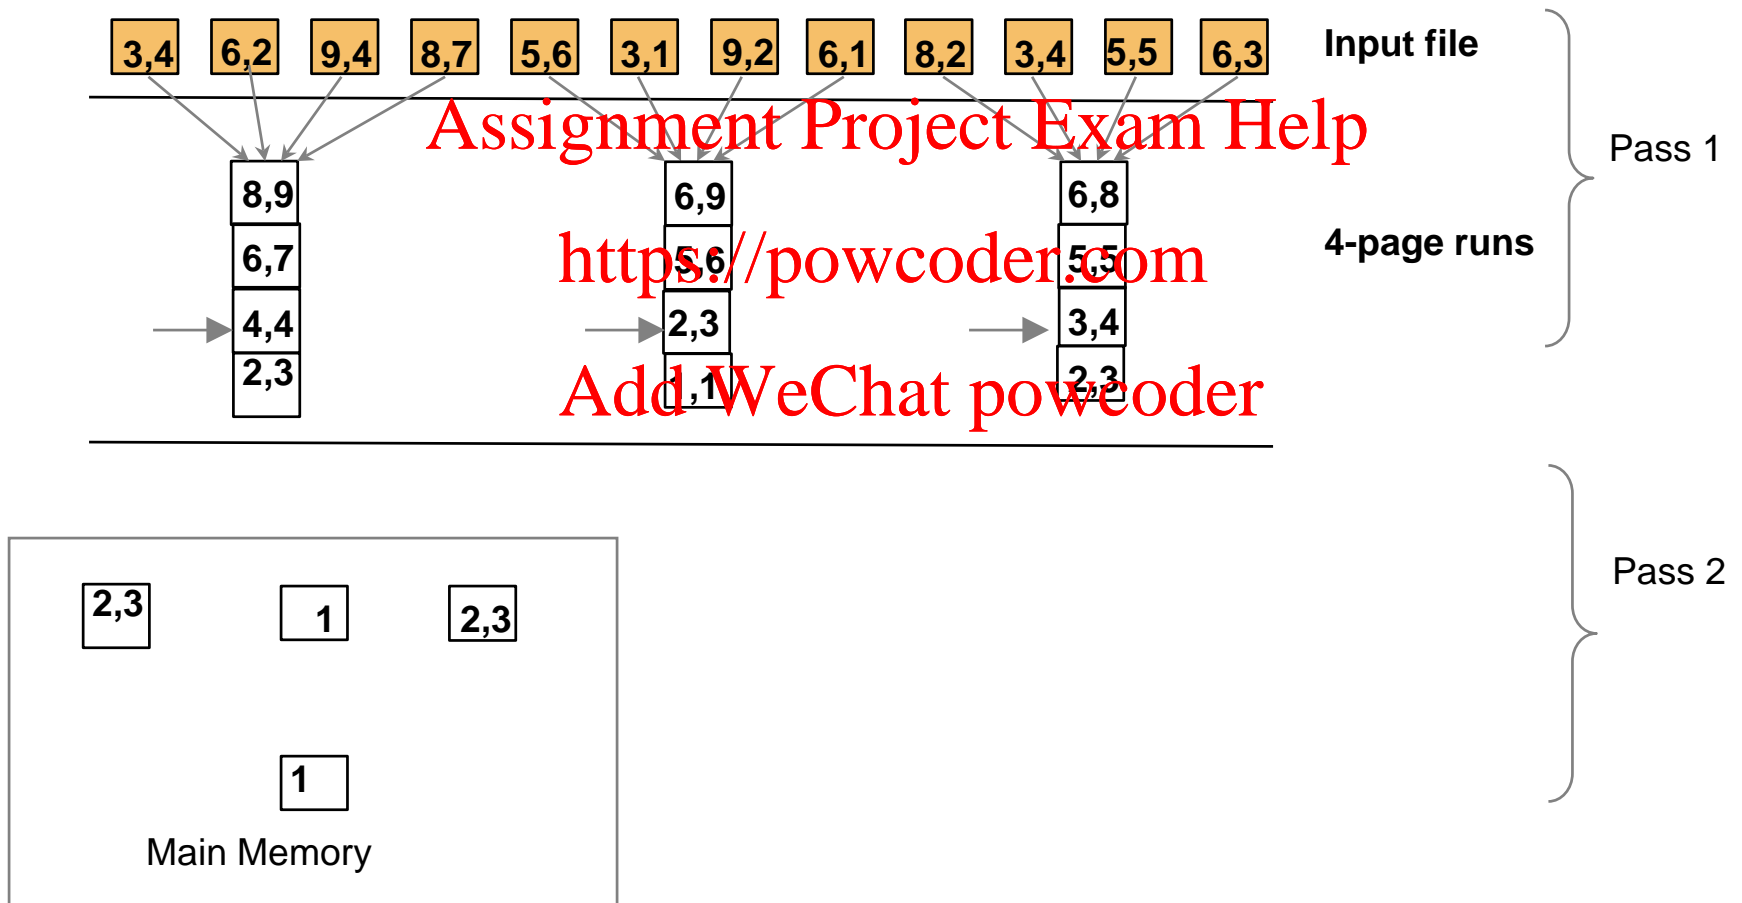
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



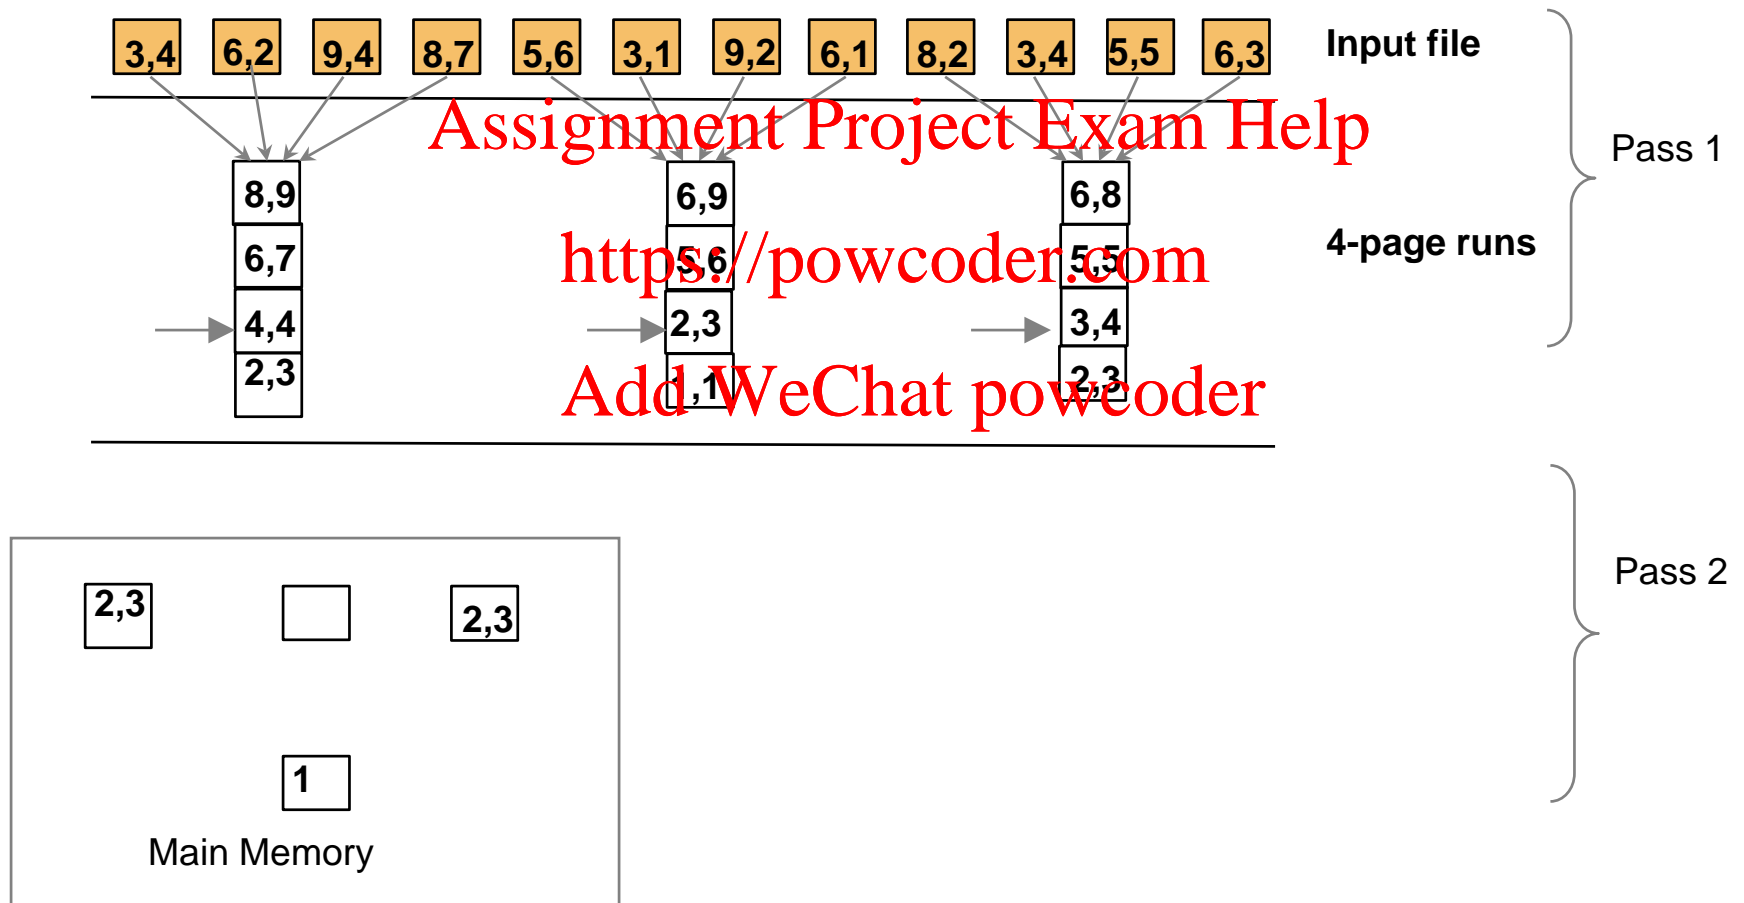
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



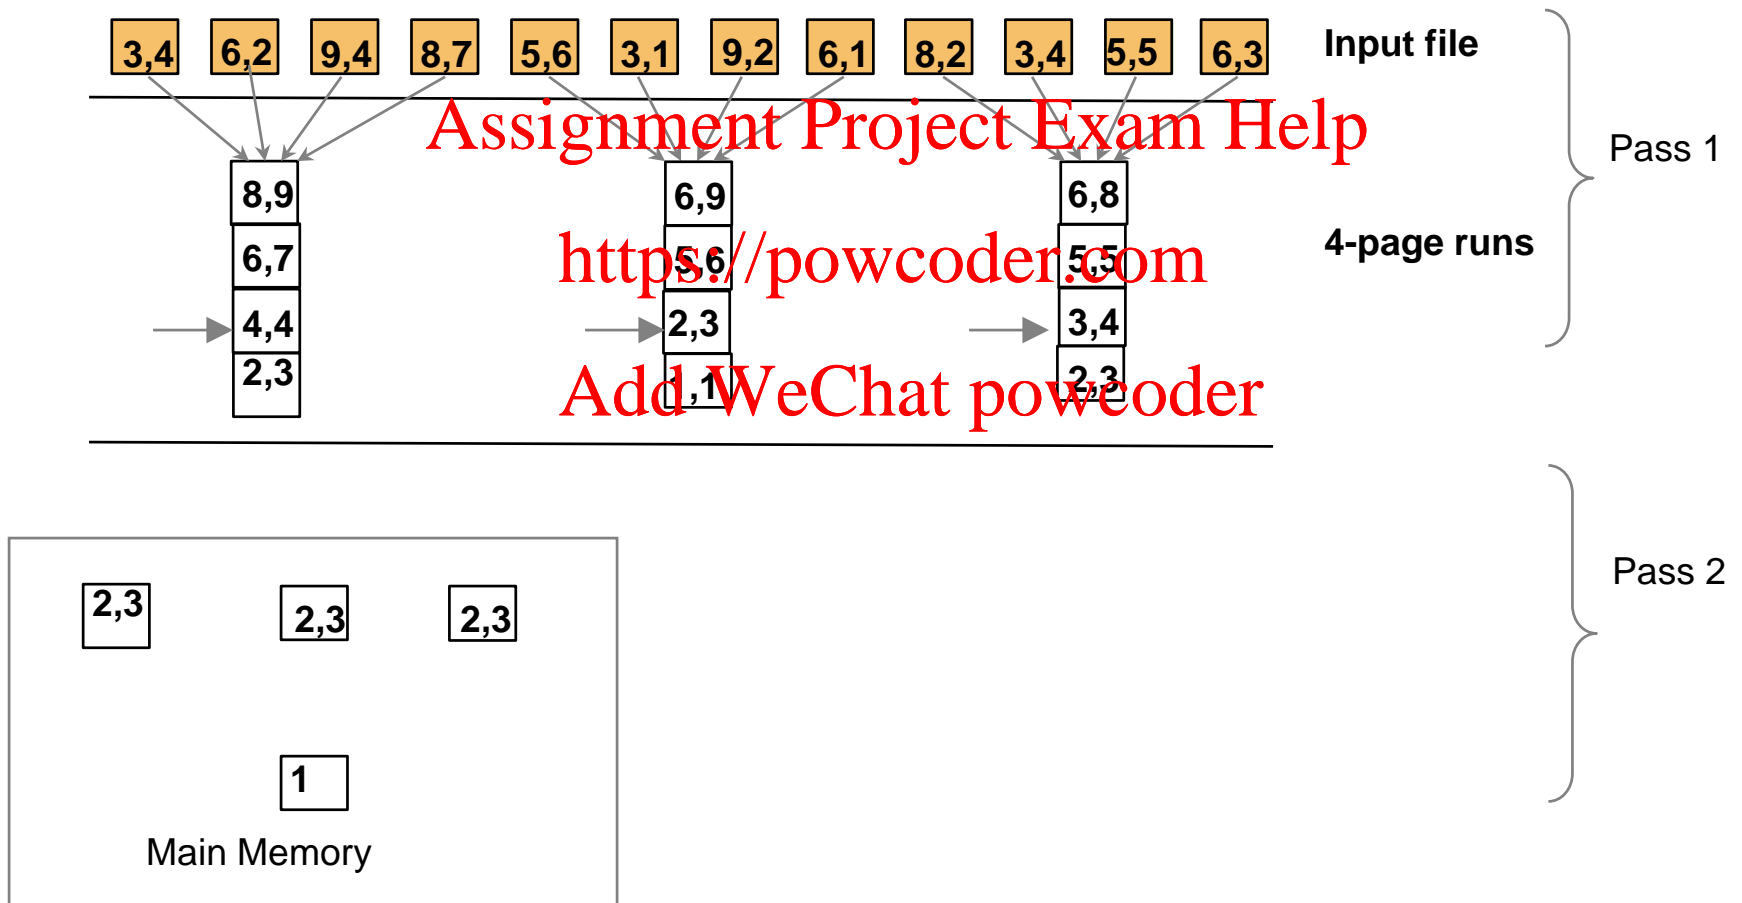
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



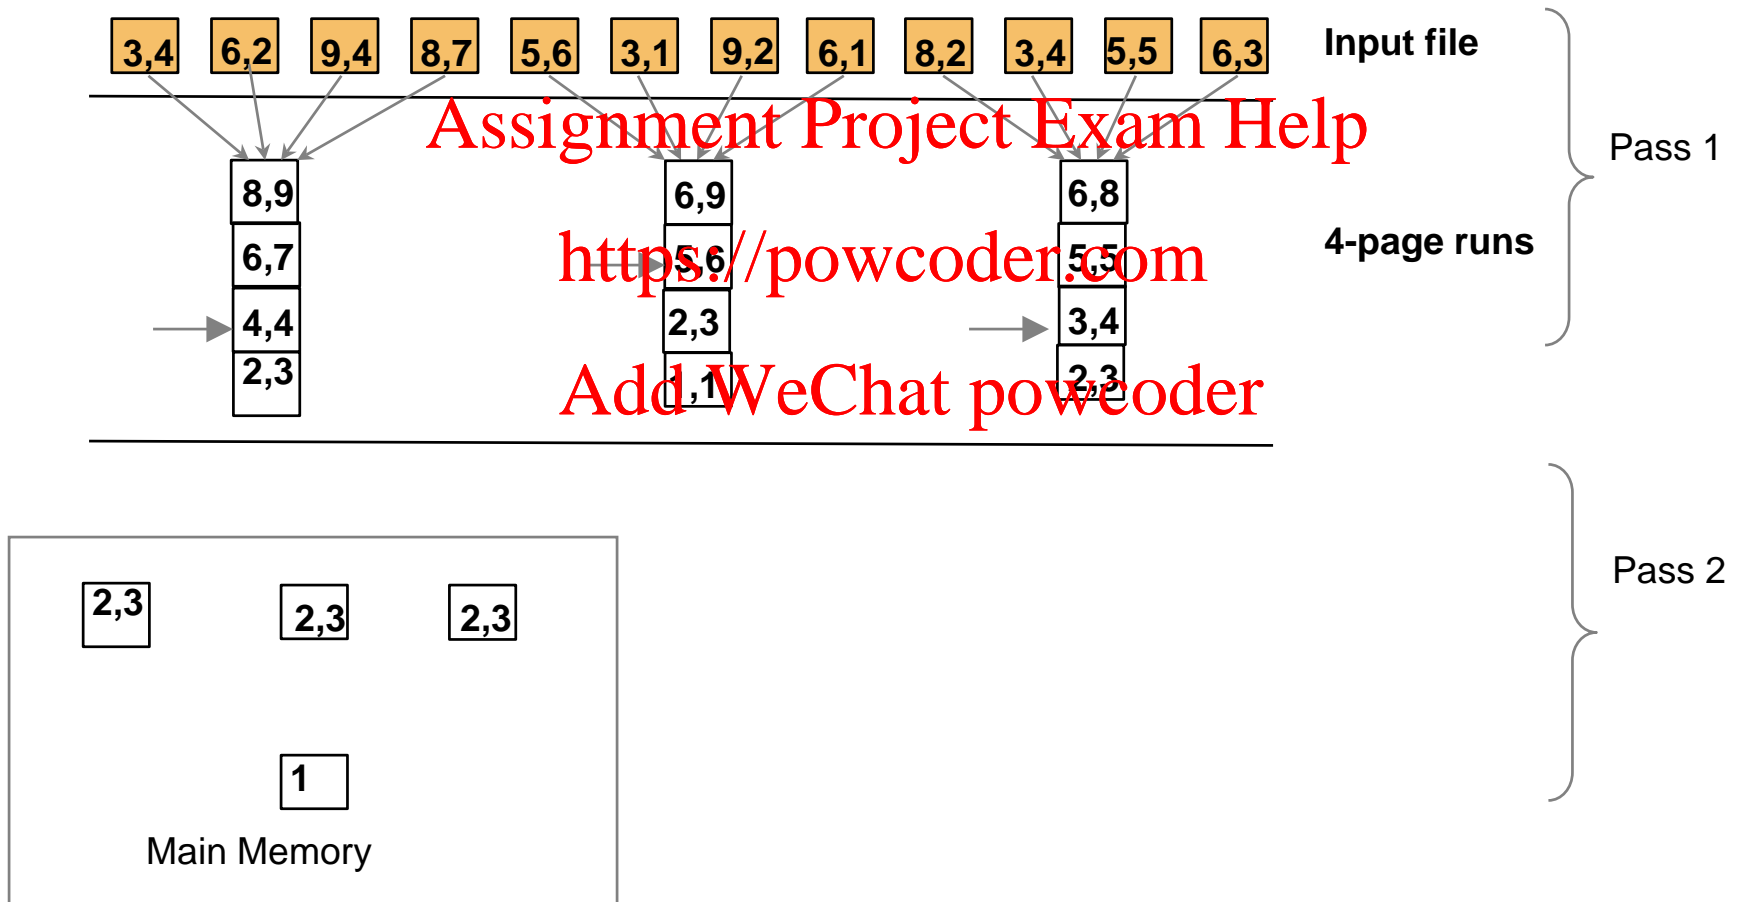
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



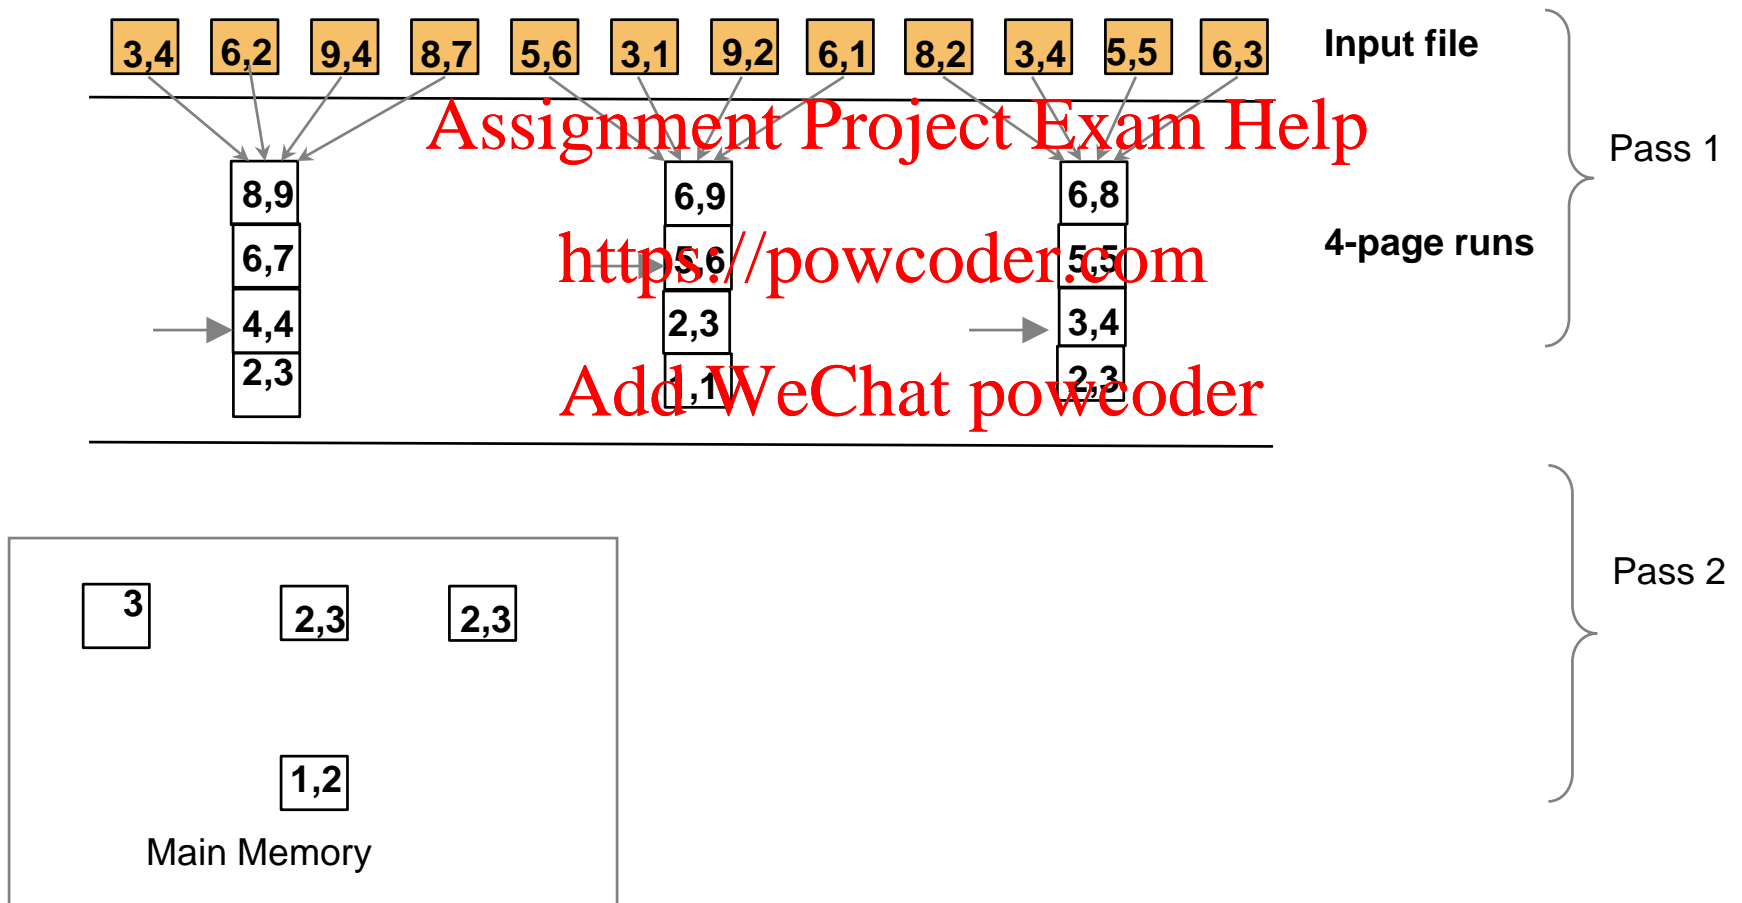
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



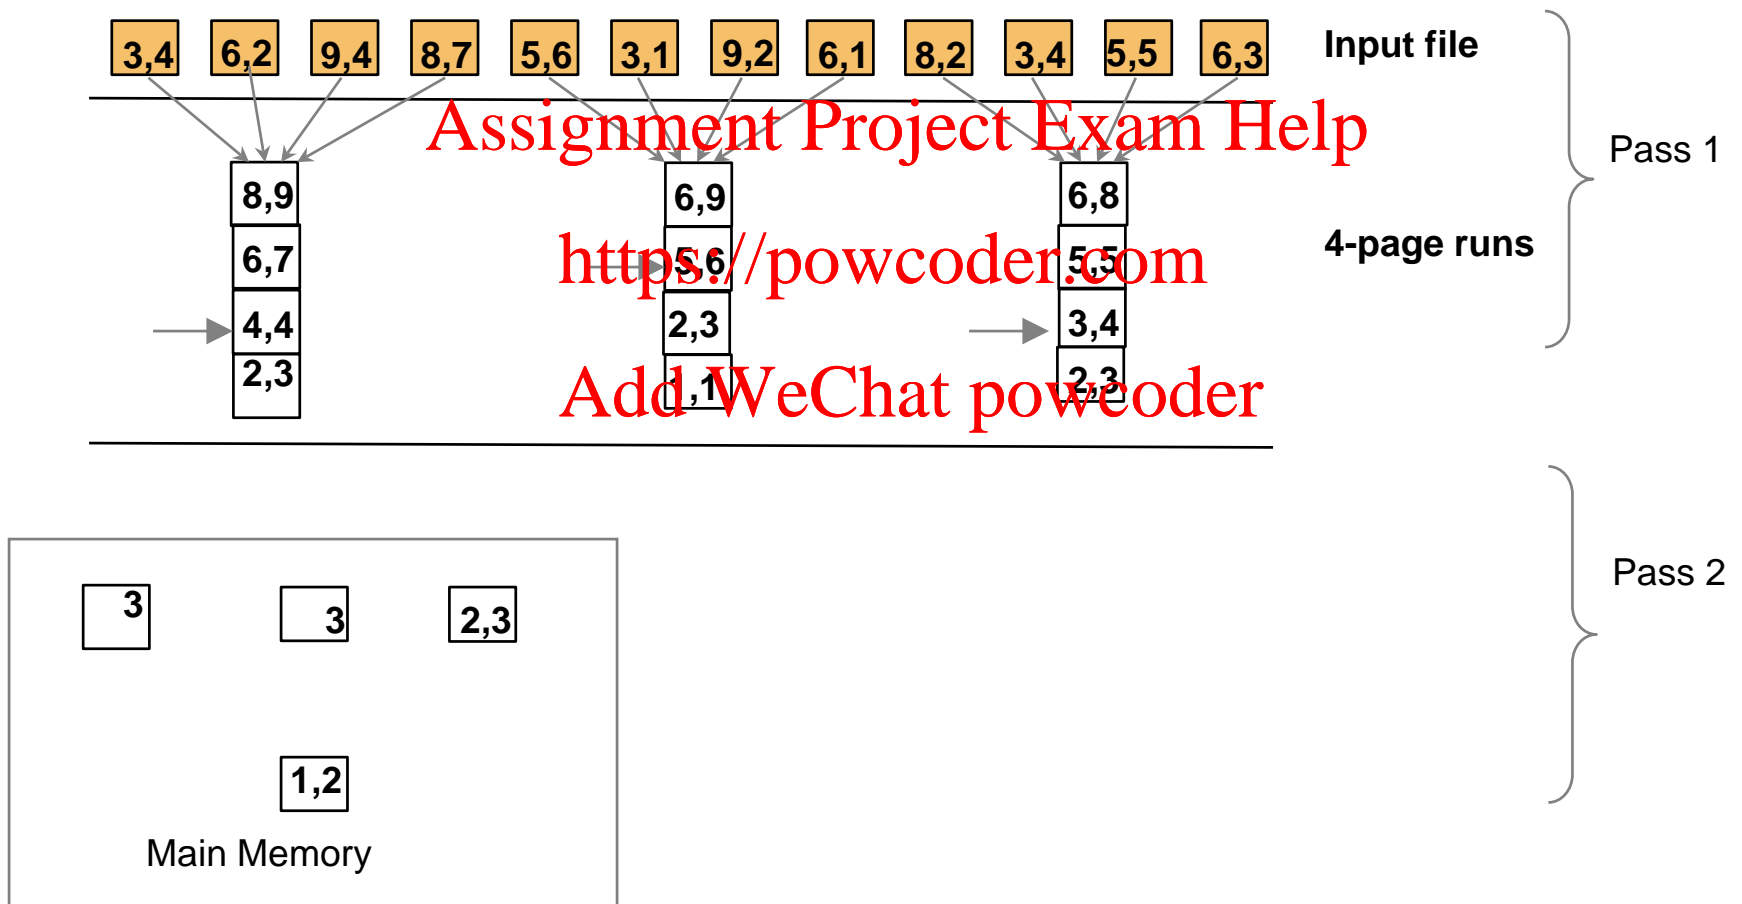
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



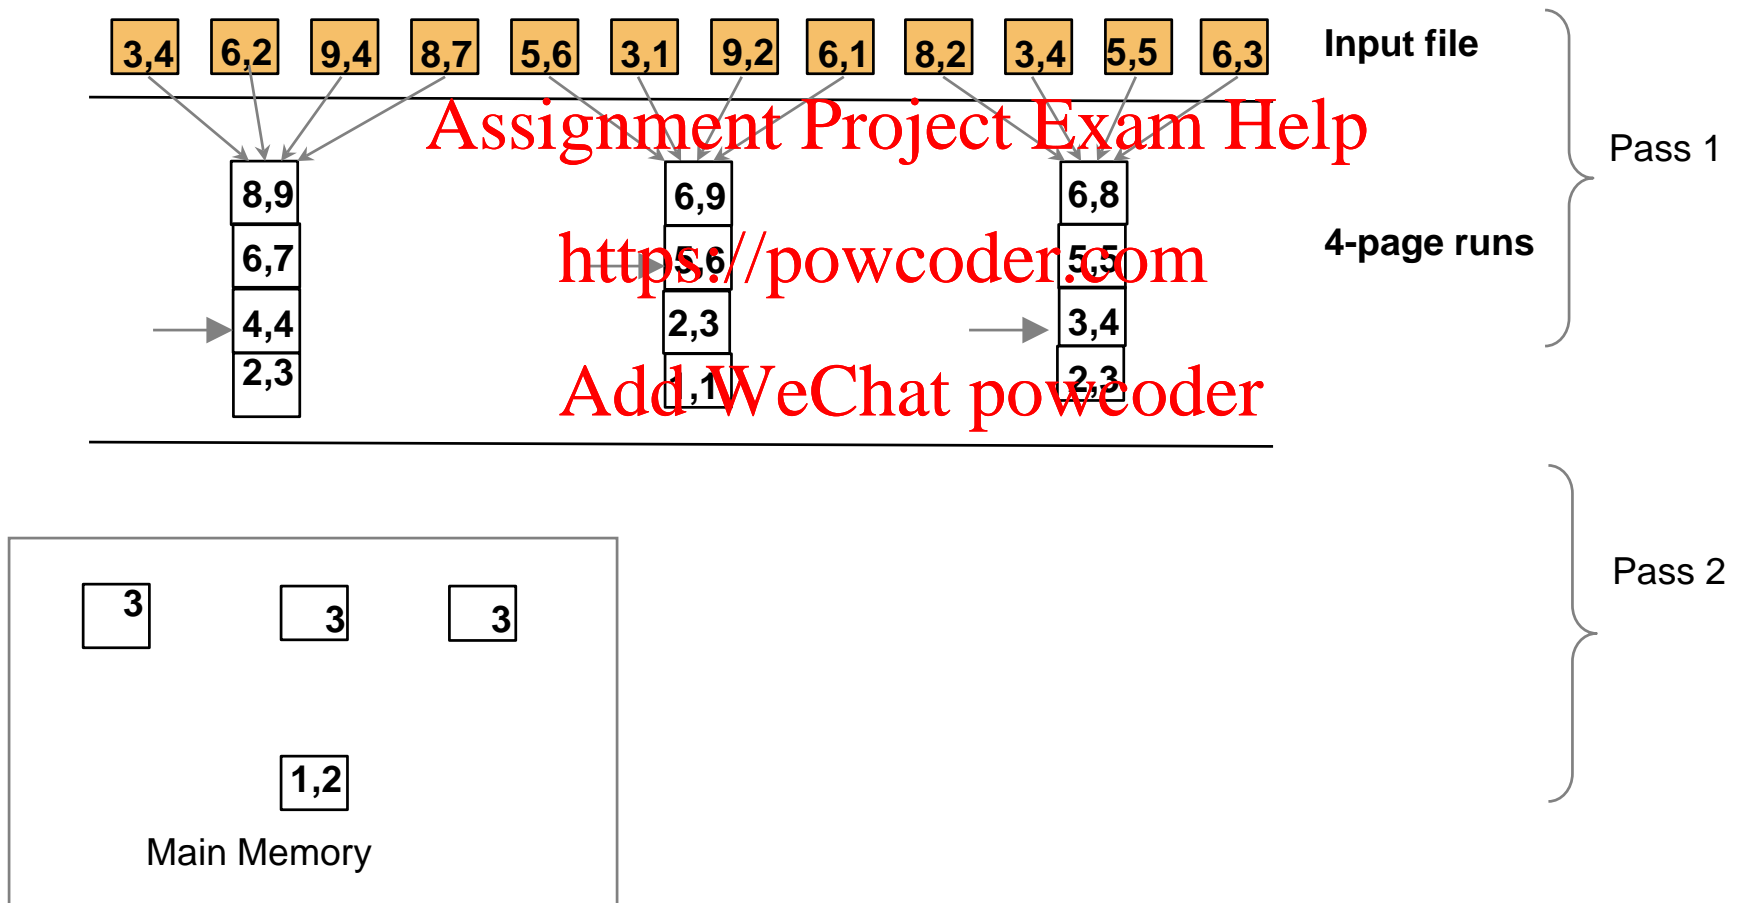
# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records



# External Merge Sort: Example

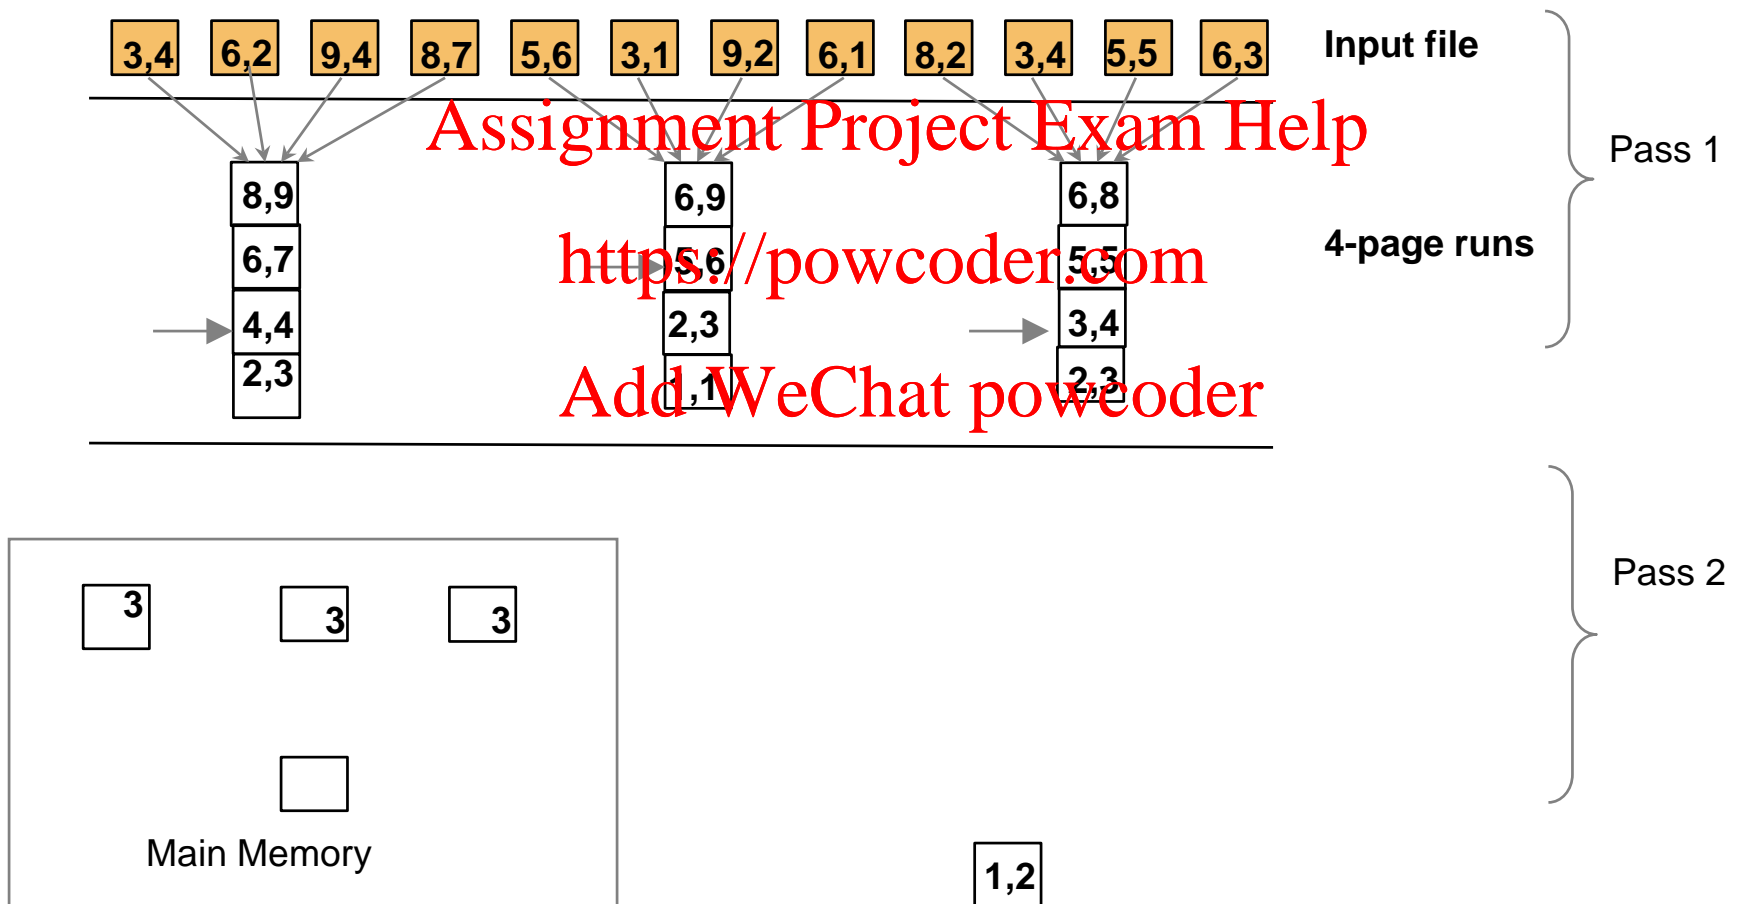
# buffer pages in memory  $B = 4$ , each page 2 records





# External Merge Sort: Example

# buffer pages in memory  $B = 4$ , each page 2 records





- Sorting with **external sort**:

- 1. Scan R, extract only the needed attributes
- 2. Sort the result set using EXTERNAL SORT
- 3. Remove adjacent duplicates

**Cost** = ReadTable + WriteProjectedPages + SortingCost + ReadProjectedPages

Read the entire table and keep only projected attributes  
Write pages with projected attributes to disk  
Sort pages with projected attributes with external sort  
Read sorted projected pages to discard adjacent duplicates

<https://powcoder.com>  
Add WeChat powcoder

**WriteProjectedPages** = NPages(R)\* PF

**PF: Projection Factor** says how much are we projecting, ratio with respect to all attributes (e.g. keeping ¼ of attributes, or 10% of all attributes)

Every time we read and write

**SortingCost** = 2\*NumPasses\*ReadProjectedPages



MELBOURNE

- **Example:** Let's say that we project  $\frac{1}{4}$  of all attributes, and let's say that we have 20 pages in memory
- $PF = \frac{1}{4} = 0.25$ ,  $NPages(R) = 1000$
- With 20 memory pages we can sort in 2 passes

Assignment Project Exam Help

**Cost** = ReadTable +  
WriteProjectedPages +  
SortingCost +  
ReadProjectedPages  
=  $1000 + 0.25 * 1000 + 2*2*250 + 250 = 2500$  (I/O)

<https://powcoder.com>

Add WeChat powcoder


MELBOURNE

- Hashing-based projection
  - 1. Scan R, extract only the **needed** attributes
  - 2. Hash data into buckets
    - Apply hash function  $h1$  to choose one of B output buffers
  - 3. Remove **adjacent** duplicates from a bucket
    - 2 tuples from different partitions guaranteed to be distinct

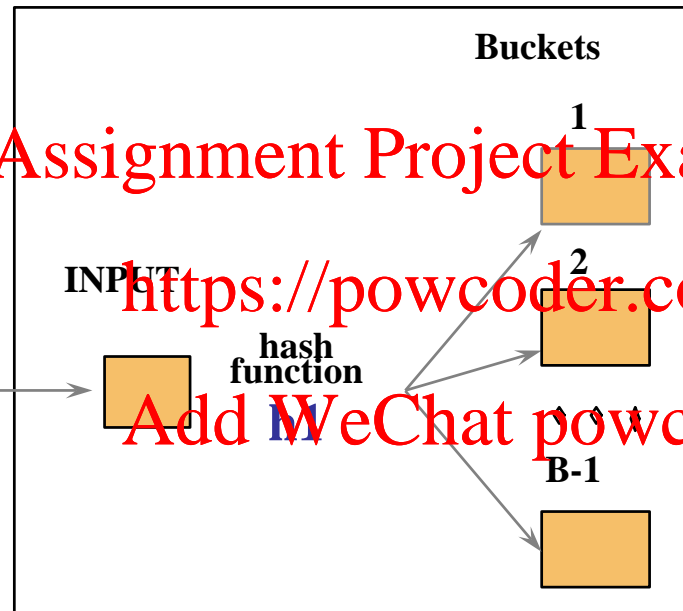
<https://powcoder.com>

Add WeChat powcoder

**Original  
Relation**



**Disk**



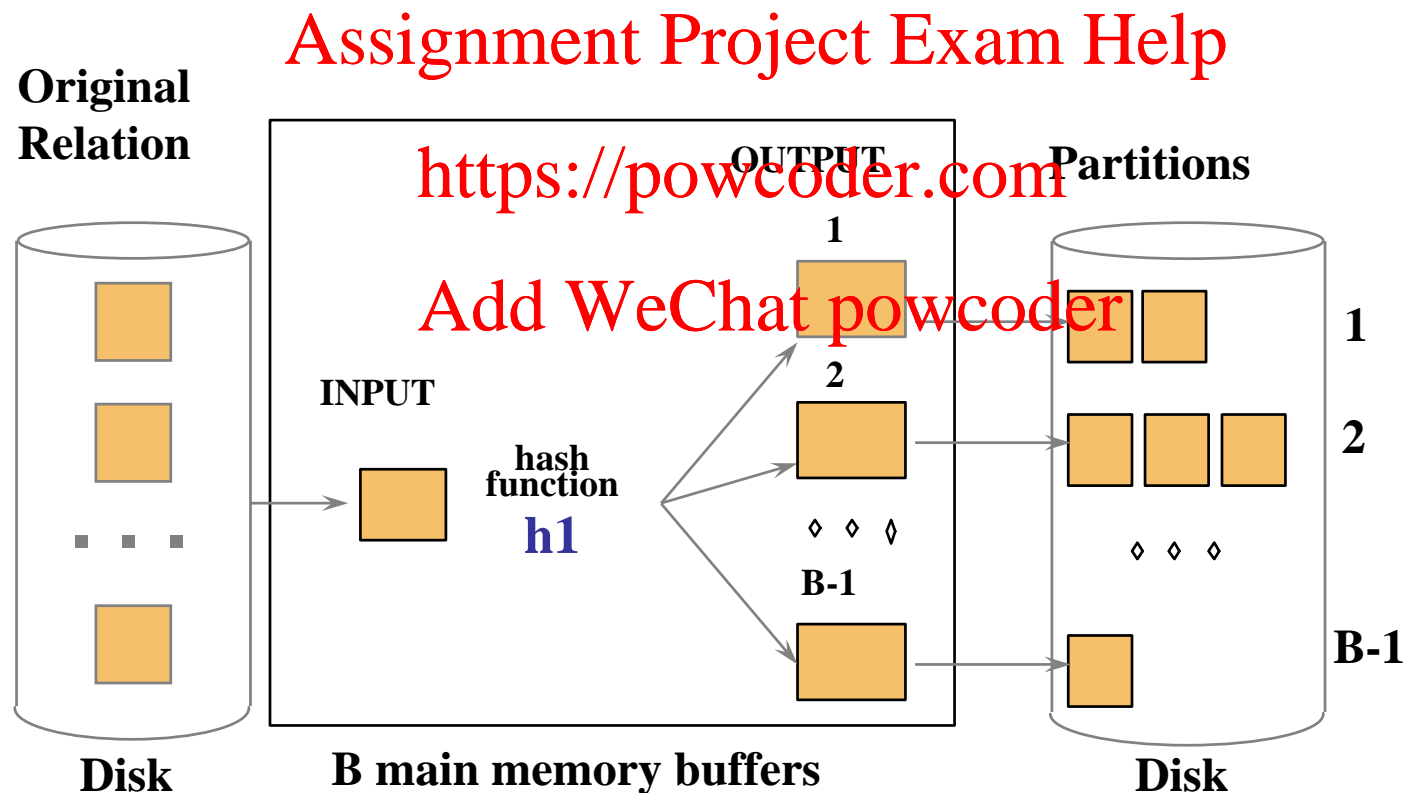
**B main memory buffers**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

1. **Partition** data into B partitions with h1 hash function
2. Load each partition, hash it with another hash function (h2) and **eliminate duplicates**



## 1. Partitioning phase:

- Read R using one input buffer
- For each tuple:
  - Discard unwanted fields
  - Apply hash function  $h1$  to choose one of B-1 output buffers
- Result is B-1 partitions (of tuples with no unwanted fields)
  - 2 tuples from different partitions guaranteed to be distinct

## 2. Duplicate elimination phase:

- For each partition
  - Read it and build an in-memory hash table
    - using hash function  $h2$  ( $\neq h1$ ) on all fields
  - while discarding duplicates
- If partition does not fit in memory
  - Apply hash-based projection algorithm recursively to this partition (we will not do this...)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



**Cost =** ReadTable + Read the entire table and project attributes  
WriteProjectedPages + Write projected pages into corresponding partitions  
ReadProjectedPages Read partitions one by one, create another hash table and discard duplicates within a bucket

## Assignment Project Exam Help

**Our example:** <https://powcoder.com>

**Cost =** ReadTable + Add WeChat powcoder  
WriteProjectedPages +  
ReadProjectedPages  
=  $1000 + 0.25 * 1000 + 250 = 1500$  (I/O)





MELBOURNE

- Understand the logic behind relational operators
- Learn alternatives for selections and projections (for now)
  - Be able to calculate the cost of alternatives
- Important for Assignment 3 as well

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

- Query Processing Part II
  - Join alternatives

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder