



INFO20003 Database Systems

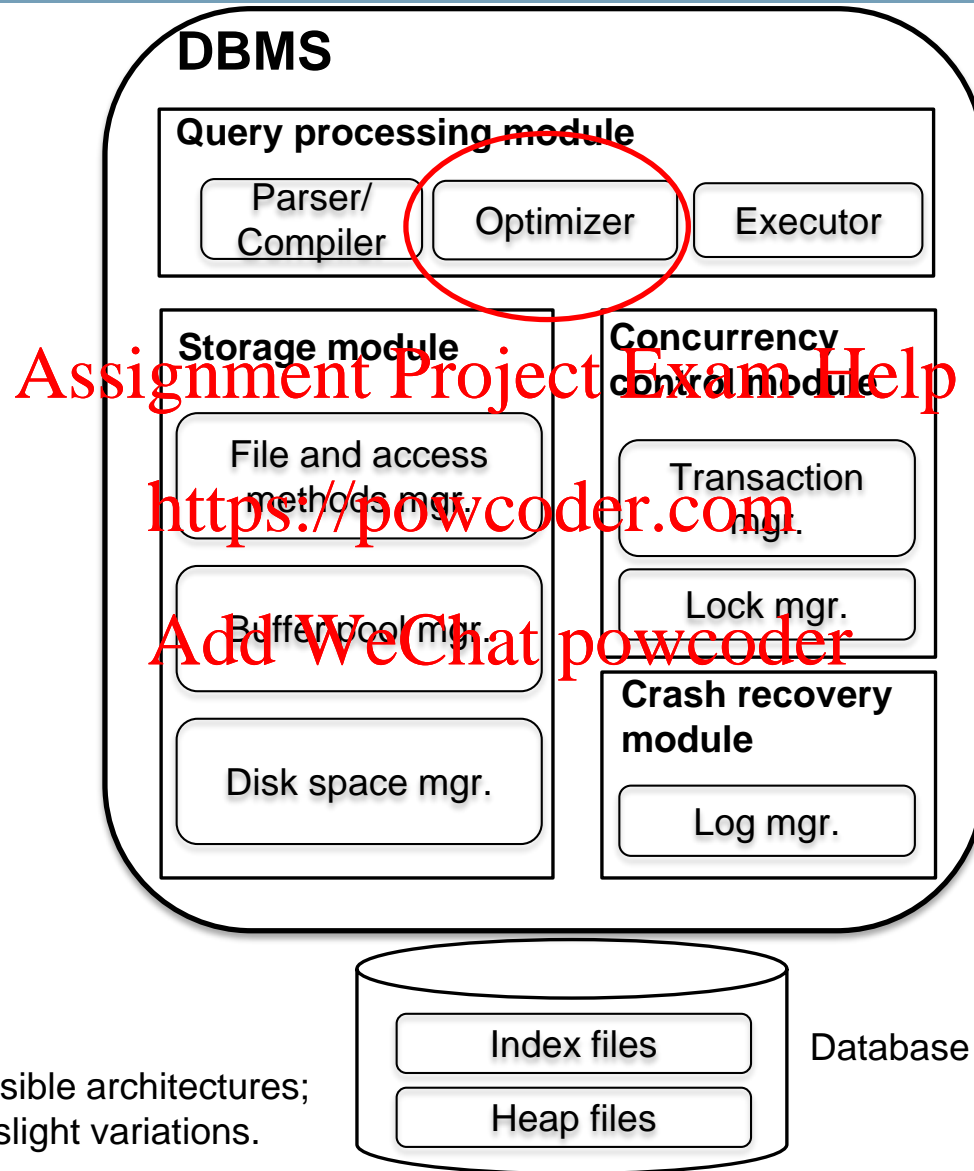
Assignment Project Exam Help

<https://powcoder.com>

Dr. Renata Borovica-Gajic
Add WeChat powcoder

Lecture 14
Query Optimization Part II

Remember this? Components of a DBMS



TODAY
Plan enumeration

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

This is one of several possible architectures; each system has its own slight variations.

Database



MELBOURNE

- When enumerating alternative plans, there are two main cases:
 - Single-relation plans
 - Multiple-relation plans (joins)
- For queries over a single relation:
 - Each *available* access path (file scan / index) is considered, and the one with the lowest estimated cost is chosen
 - Heap scan is always one alternative
 - Each index can be another alternative (if matching selection predicates)
 - Other operations can be performed on top of access paths, but they typically don't incur additional cost since they are done *on the fly* (e.g. projections, additional non-matching predicates)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

1. Sequential (heap) scan of data file:

$$\text{Cost} = \text{NPages}(\mathbf{R})$$

2. Index selection over a primary key (just a single tuple):

$$\text{Cost}(\mathbf{B+Tree}) = \text{Height}(\mathbf{I}) + 1$$

Height is the index height

$$\text{Cost}(\text{HashIndex}) = \text{ProbeCost}(\mathbf{I}) + 1, \text{ProbeCost}(\mathbf{I}) \sim 1.2$$

<https://powcoder.com>

3. Clustered index matching one or more predicates:

$$\text{Cost}(\mathbf{B+Tree}) = (\text{NPages}(\mathbf{I}) + \text{NPages}(\mathbf{R})) * \prod_{i=1..n} RF_i$$

$$\text{Cost}(\text{HashIndex}) = \text{NPages}(\mathbf{R}) * \prod_{i=1..n} RF_i * 2.2$$

4. Non-clustered index matching one or more predicates:

$$\text{Cost}(\mathbf{B+Tree}) = (\text{NPages}(\mathbf{I}) + \text{NTuples}(\mathbf{R})) * \prod_{i=1..n} RF_i$$

$$\text{Cost}(\text{HashIndex}) = \text{NTuples}(\mathbf{R}) * \prod_{i=1..n} RF_i * 2.2$$



Let's say that Sailors(S) has 500 pages, 40000 tuples, NKeys(rating) = 10

```
SELECT S.sid FROM Sailors S WHERE S.rating=8
```

• Result size = $(1/NKeys(rating)) * NTuples(S) = (1/10) * 40000 = 4000$ tuples

1. If we have $I(rating)$, NPages(I) = 50:

– Clustered index:

Cost = $(1/NKeys(rating)) * (NPages(I) + NPages(S)) = (1/10) * (50 + 500) = 55$ I/O

– Unclustered index:

Cost = $(1/NKeys(rating)) * (NPages(I) + NTuples(S)) = (1/10) * (50 + 40000) = 4005$ I/O

2. If we have an $I(sid)$, NPages(I) = 50:

– Cost = ?, Result size = ?

– Would have to retrieve all tuples/pages. With a clustered index, the cost is 50+500, with unclustered index, 50+40000

3. Doing a file scan:

– Cost = NPages(S) = 500

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Steps:

1. Select **order** of relations
 - E.g. SxRxB, or SxBxR or RxSxB...
 - maximum possible orderings = $N!$
2. For each join, select join algorithm
 - E.g. Hash join, Sort-Merge Join...
3. For each input relation, select access method
 - Heap Scan, or various index alternatives

Q: How many plans are there for a query over N relations?

Back-of-envelope calculation:

- With 3 join algorithms, I indexes per relation:
plans $\approx [N!] * [3^{(N-1)}] * [(I + 1)^N]$
- Suppose $N = 3$, $I = 2$: # plans $\approx 3! * 3^2 * 3^3 = 1458$ plans
- This is just for illustration – you don't need to remember this





Queries Over Multiple Relations

- As number of joins increases, number of alternative plans grows rapidly → *need to restrict search space*
- Fundamental decision in System R (first DBMS): **only left-deep join trees** are considered
 - Left-deep trees allow us to generate all *fully pipelined* plans
 - Intermediate results are not written to temporary files

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



```
SELECT S.sname, B.bname, R.day  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid
```

- Let's assume:

- Two join algorithms to choose from
 - Hash-Join
 - NL-Join (page-oriented)
- Clustered B+Tree index: $I(R.sid); NPages(I) = 50$
- No other indexes
- S: $NPages(S) = 500, NTuplesPerPage(S) = 80$
- R: $NPages(R) = 1000, NTuplesPerPage(R) = 100$
- B: $NPages(B) = 10$
- 100 R \bowtie S tuples fit on a page

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Candidate Plans

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```






1. Enumerate relation orderings:



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



* Prune plans with cross-products immediately!

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```





2. Enumerate join algorithm choices:



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



+ do the same
for other plans

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

3. Enumerate access method choices:





+ do same for
other plans


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Now estimate the cost of each plan



```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

S: NPages(S) = 500, NTuplesPerPage(S) = 80
 R: NPages(R) = 1000, NTuplesPerPage(R) = 100

B: NPages(B) = 10
 100 R \bowtie S tuples fit on a page
 All 3 relations are Heap Scan

Assignment Project Exam Help

<https://powcoder.com>

Calculating cost:

SxR

Cost (SxR) = 500 + 500 * 1000 = 500500

(SxR)xB

Result size (SxR) = 40000 * 100000 * ~~1~~ / 40000 = 100000 tuples = 1000 pages


Cost(xB) = ~~1000~~ + 1000 * 10 = 10000

Already read – left deep plans apply pipelining

Total Cost = 500 + 500 * 1000 + 1000 * 10 = 510500 I/O

Add WeChat powcoder

Now estimate the cost of each plan



S: NPages(S) = 500, NTuplesPerPage(S) = 80
R: NPages(R) = 1000, NTuplesPerPage(R) = 100
B: NPages(B) = 10
100 R \bowtie S tuples fit on a page
All 3 relations are Heap Scan

Assignment Project Exam Help

Calculating cost:

<https://powcoder.com>

SxR

Cost (SxR) = 500 + 500*1000 = 500500

Add WeChat powcoder

(SxR)xB

Result size (SxR) = 100000*40000 * $\frac{1}{40000}$ = 100000 tuples = 1000 pages


Cost(xB) = $\frac{3}{2} \times 1000 + 3 \times 10 = 2 \times 1000 + 3 \times 10 = 2030$

Already read once – left deep plans apply pipelining

Total Cost = 500 + 500*1000 + 2*1000 + 3*10 = 502530 I/O



Plan 3:




S: $NPAGES(S) = 500$, $NTuplesPerPage(S) = 80$
 R: $NPAGES(R) = 1000$, $NTuplesPerPage(R) = 100$
 B: $NPAGES(B) = 10$
 $100 R \bowtie S$ tuples fit on a page
 All 3 relations are Heap Scan

Assignment Project Exam Help

<https://powcoder.com>

Plan 4:



Calculating cost:
 Add WeChat powcoder
 Cost (P3) = ?

Cost (P4) = ?



S: $NPAGES(S) = 500$, $NTuplesPerPage(S) = 80$


R: $NPAGES(R) = 1000$, $NTuplesPerPage(R) = 100$

B: $NPAGES(B) = 10$, $NTuplesPerPage(B) = 10$


SMJ : 2 passes, RxB: 10 tuples per page

I(S.sid); $NPAGES(I) = 50$


1)



2)



3)



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

- Understand plan enumeration and cost various plans
- Important for Assignment 3 as well

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



MELBOURNE

- Normalization

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder