



# INFO20003 Database Systems

Assignment Project Exam Help

<https://powcoder.com>

Dr Renata Borovica-Gajic\*

Add WeChat powcoder

Lecture 20

Distributed Databases

Week 10

*slides adopted  
from David Eccles*

- What is a distributed database?
- Why are they used, and how they work
- Pros and cons of different approaches

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



distributed database



replicated database

material in this lecture is drawn from Hoffer et al. (2013) *Modern Database Management* 11<sup>th</sup> edition, chapter 12, available online at [http://wps.prenhall.com/bp\\_hoffer\\_mdm\\_11/230/58943/15089539.cw/index.html](http://wps.prenhall.com/bp_hoffer_mdm_11/230/58943/15089539.cw/index.html)  
pictures on this page are from Gillenson (2005) *Fundamentals of Database Management Systems*

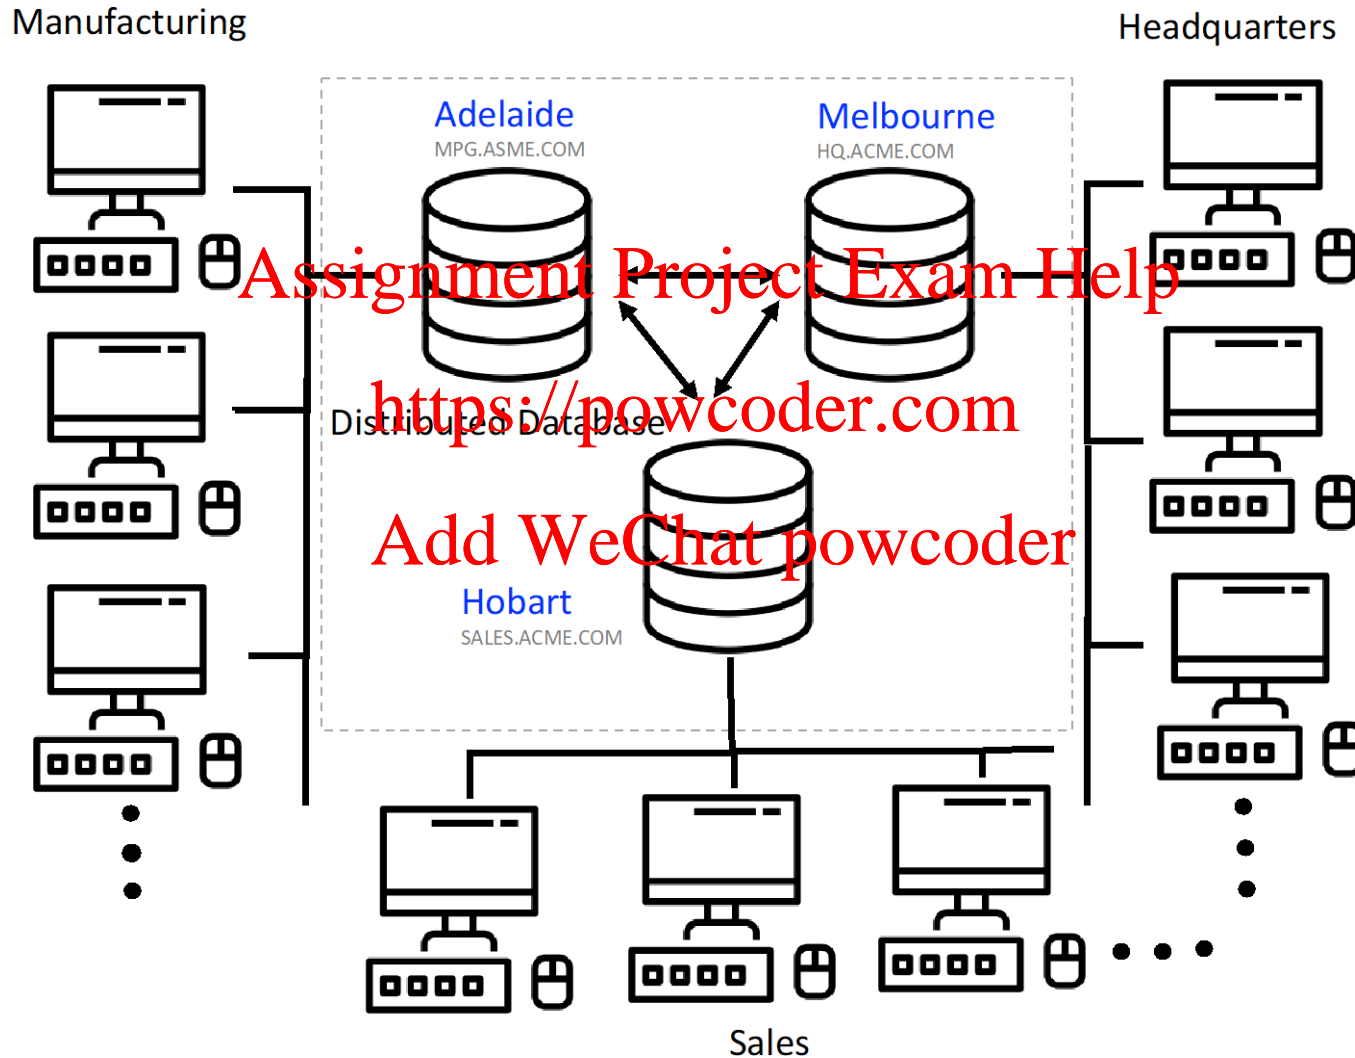
- Distributed Database
  - a single logical database physically spread across multiple computers in multiple locations that are connected by a data communications link
  - *appears to users as though it is one database*
- Decentralized Database
  - a collection of independent databases which are not networked together as one logical database
  - *appears to users as though many databases*
- We are concerned with *distributed* databases

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Example – distributed database




# Advantages of distributed DBMS

- Good fit for geographically distributed organizations / users
  - Utilize the internet
- Data located near site with greatest demand
  - E.g. ESPN Weekend Sports Scores




- Faster data access (to local data)
- Faster data processing
  - Workload split amongst physical servers



Vertical scaling

VS.



Horizontal scaling

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


- Allows modular growth
  - add new servers as load increases (horizontal scalability)



- Increased reliability and availability
  - less danger of a single point of failure (SPOF), IF data is replicated



- Supports database recovery
  - When data is replicated across multiple sites



# Disadvantages of distributed DBMS

- Complexity of management and control
  - Database or/and application must stitch together data across sites
    - Who and where is the current version of the record (row & column)?
    - Who is waiting to update that information and where are they?
    - How does the log get displayed to the web & application server?
- Data integrity
  - Additional exposure to improper updating
    - If two users in two locations update the record at the exact same time who decides which statement should “win”?
  - Solution: Transaction Manager or Master-slave design
- Security
  - Many server sites -> higher chance of breach
    - Multiple access sites require protection including network and storage infrastructure from both cyber & physical attacks

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- Lack of standards
  - Different Relational DDBMS vendors use different protocols
- Increased training & maintenance costs
  - More complex IT infrastructure
  - Increased Disk storage (\$)
  - Fast intra and inter network infrastructure (\$\$\$)
  - Clustering software (\$\$\$\$)
  - Network Speed (\$\$\$\$\$)
- Increased storage requirements
  - Replication model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





- Location transparency
  - a user does not need to know where particular data are stored
- Local autonomy
  - a node can continue to function for local users if connectivity to the network is lost

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


- A user (or program) accessing data do not need to know the location of the data in the network of DBMS's
- Requests to retrieve or update data from any site are automatically forwarded by the system to the site or sites related to the processing request
- A single query can join data from tables in multiple sites

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
SELECT hometeam, homescore, awayteam, awayscore  
FROM results INNER JOIN codes  
ON results.codeid = codes.codeid  
WHERE sportscore in ('NFL', 'Hurling', 'EPL');
```



- Being able to operate locally when connections to other databases fail
- Users can administer their local database
  - control local data (e.g Hurling results)
  - administer security
  - log transactions
  - recover when local failures occur
  - provide full access to local data

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Functions of a distributed DBMS

- Locate data with a distributed catalog (meta data)
- Determine location from which to retrieve data and process query components
- DBMS translation between nodes with different local DBMSs (using middleware)
- Data consistency (via multiphase commit protocols)
- Global primary key control
- Scalability
- Security, concurrency, query optimization, failure recovery

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


# Distribution options

- When distributing data around world – the data can be *partitioned* or *replicated*.
- Data replication** is a process of *duplicating* data to different nodes.
- Data partitioning** is the process of *partitioning* data into subsets that are shipped to different nodes.
- Many real-life systems use a combination of two (e.g. partition data and keep some replicas around – usually 3)

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



# Replication - advantages

- High reliability due to redundant copies of data
- Fast access to data at the location where it is most accessed
- May avoid complicated distributed integrity routines
  - Replicated data is refreshed at scheduled intervals
- Decoupled nodes don't affect data availability
  - Transactions proceed even if some nodes are down
- Reduced network traffic at prime time
  - If updates can be delayed
- This is currently popular as a way of achieving high availability for global systems
  - Most SQL & NoSQL databases offer replication



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

- Need more storage space
  - Each server stores a *copy* of the row
- Data Integrity:
  - High tolerance for out-of-date data may be required
  - Updates may cause performance problems for busy nodes
  - Retrieve incorrect data if updates have not arrived


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Centralised Database  
One database in one server  
(1 copy of data)



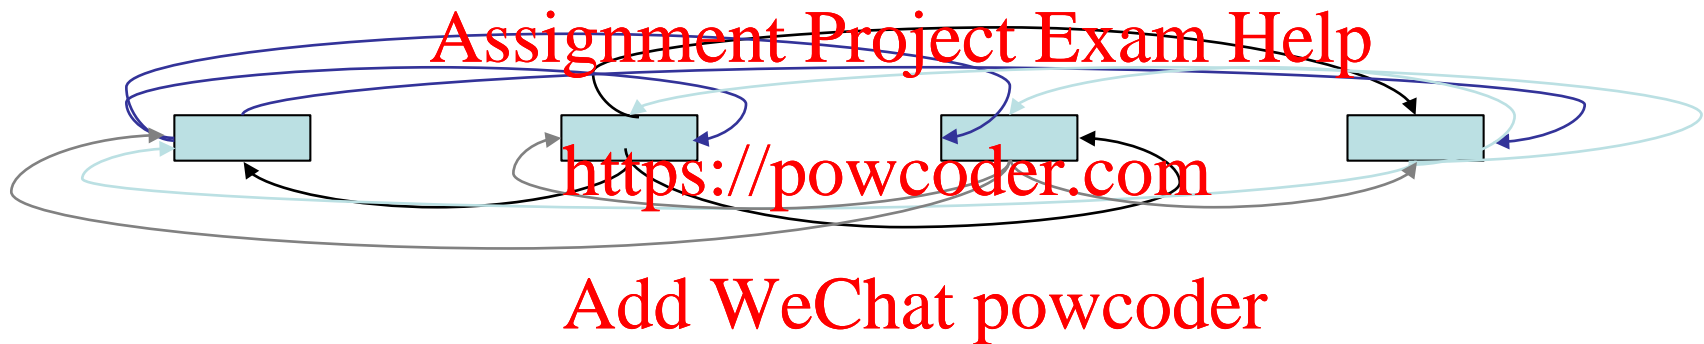
Distributed (Replicated) Database  
One database in 4 physical servers  
(4 copies of data)



Data Size

# Replication - disadvantages

- Takes time for update operations
  - High tolerance for out-of-date data may be required
  - Updates may cause performance problems for busy nodes



- Network communication capabilities
  - Updates can place heavy demand on telecommunications/networks
  - High speed networks are expensive (\$\$\$\$\$)



- Split data into chunks, store chunks in different nodes
- A chunk can be a set of rows or columns
- Thus, two types of partitioning: horizontal & vertical

Assignment Project Exam Help

- **Horizontal partitioning**

- Table rows distributed across nodes (sides)

- **Vertical partitioning**

- Table columns distributed across nodes (sides)

<https://powcoder.com>

Add WeChat powcoder

# Horizontal partitioning

- Different rows of a table at different sites
- Advantages
  - data stored close to where it is used
    - efficiency
  - local access optimization
    - better performance
  - only relevant data is stored locally
    - security
  - unions across partitions
    - ease of query
- Disadvantages
  - accessing data across partitions
    - inconsistent access speed
  - no data replication
    - backup vulnerability (SPOF)

Team table

ID	Team	City	Code	Region	League
1	Arsenal	London	Football	Europe	EPL
2	Jets	NYC	Grid Iron	Americas	NFL
3	Carlton FC	Melbourne	Aussie Rules	APAC	AFL
4	Racing92	Paris	Rugby	Europe	Top14
5	Yankees	NYC	Baseball	Americas	MLB
6	Swifts	Sydney	Netball	APAC	ANZ

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Example horizontal partitioning

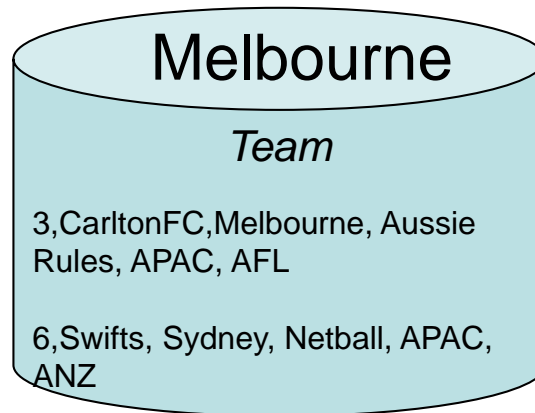
ID	Team	City	Code	Region	League
1	Arsenal	London	Football	Europe	EPL
2	Jets	NYC	Grid Iron	Americas	NFL
3	Carlton FC	Melbourne	Aussie Rules	APAC	AFL
4	Racing92	Paris	Rugby	Europe	Top14
5	Yankees	NYC	Baseball	Americas	MLB
6	Swifts	Sydney	Netball	APAC	ANZ

Assignment Project Exam Help

<https://powcoder.com>

Horizontal Partitioning based on Region

Add WeChat powcoder






# Vertical partitioning

- Different columns of a table at different sites
- Advantages and disadvantages are the same as for horizontal partitioning, *except*
  - combining data across partitions is more difficult because it requires joins (instead of unions)

Assignment Project Exam Help

<https://powcoder.com>

Player table

ID	Firstname	Lastname	Team	League	Photo	Biography
110	Luc	Eucalon	4	Top14		Ipso locum
120	Vasil	Kakokan	4	Top14		Ipso locum est
130	Donacca	Ryan	4	Top14	<null>	
210	Edwin	Maka	4	Top14		

Add WeChat powcoder

# Example vertical partitioning

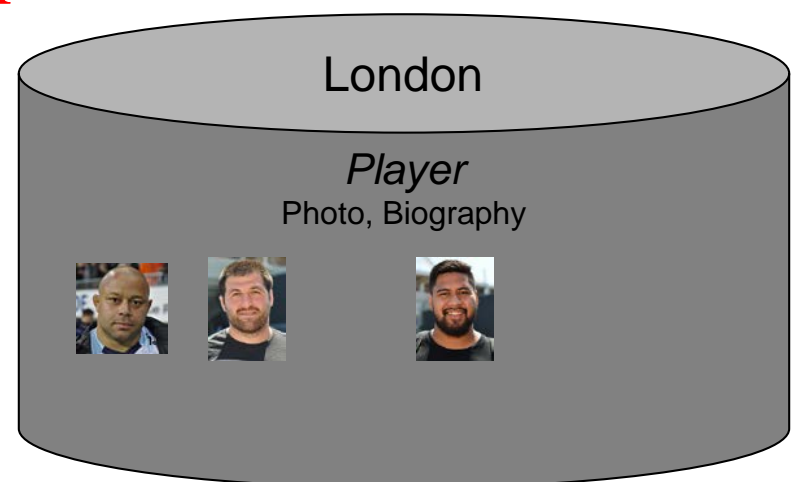
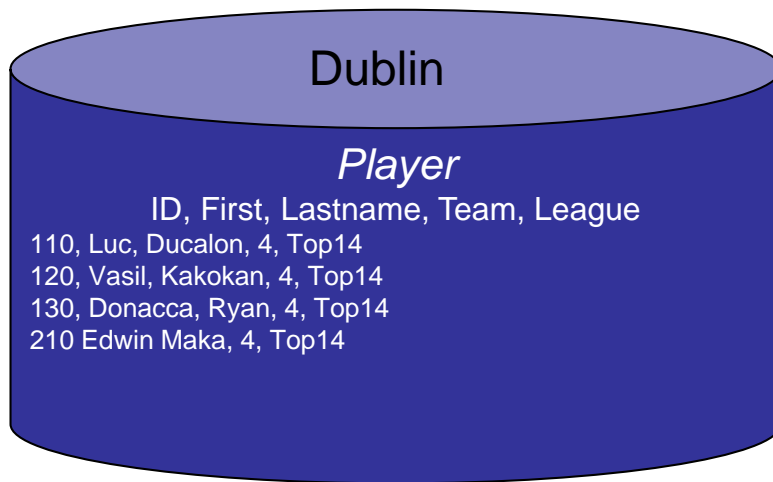
ID	Firstname	Lastname	Team	League	Photo	Biography
110	Luc	Ducalon	4	Top14		Ipsa locum
120	Vasil	Kakokan	4	Top14		Ipsa locum est
130	Donacca	Ryan	4	Top14		
210	Edwin	Maka	4	Top14		

Assignment Project Exam Help

<https://powcoder.com>

Vertical Partitioning based on column requirements

Add WeChat powcoder



- Trade-offs
  - Availability vs Consistency
    - The CAP theorem says we need to decide whether to make data always available OR always consistent
  - Synchronous vs Asynchronous updates
    - Are changes immediately visible everywhere (great BUT expensive) or later propagated (less expensive faster, but seeing stale data)?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# The CAP theorem

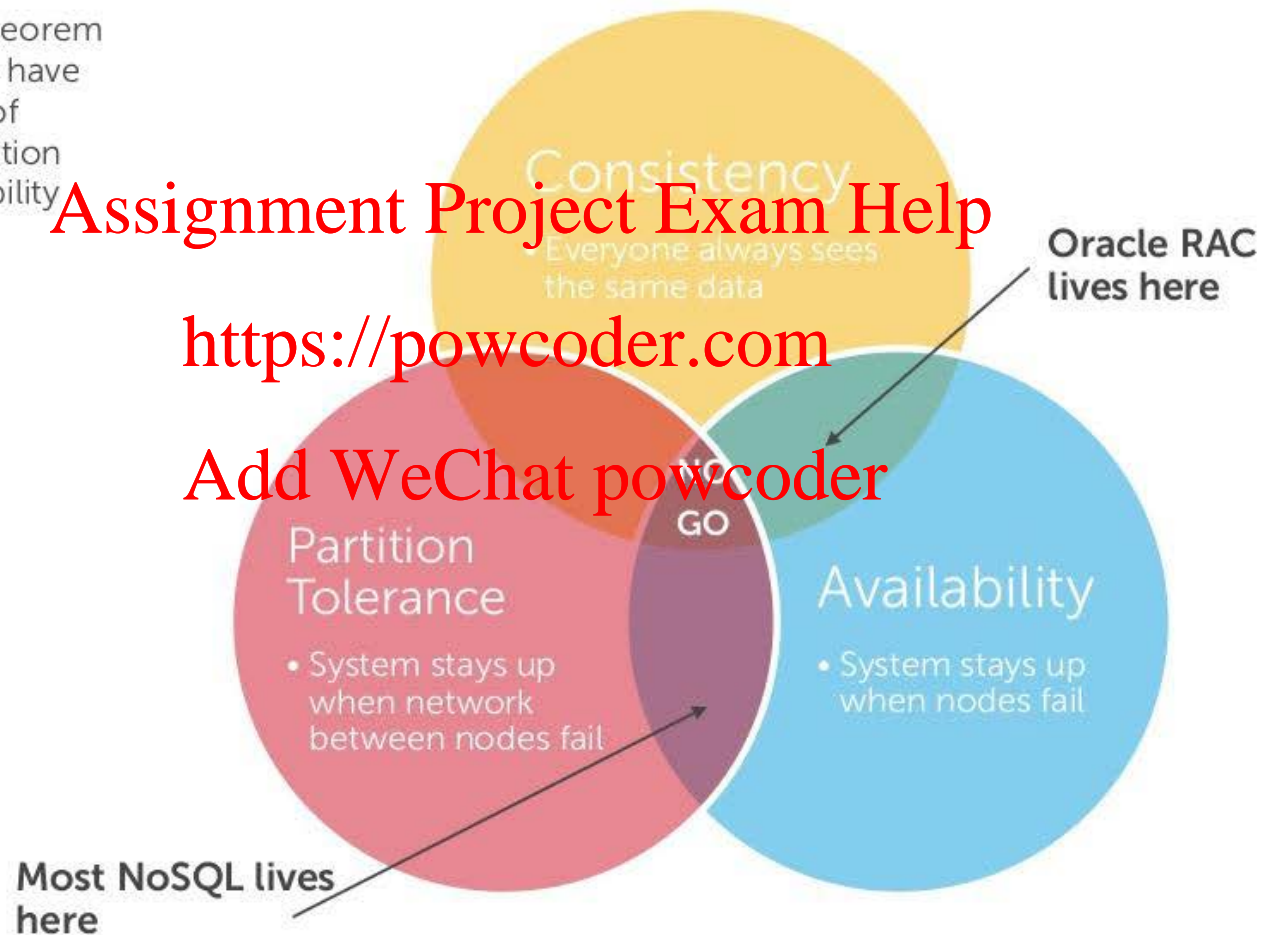
CAP Theorem says something has to give

- CAP (Brewer's) Theorem says you can only have two out of three of Consistency, Partition Tolerance, Availability

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Synchronous updates

- Data is continuously kept up to date
  - users anywhere in the world can access data and get the same answer
- If any copy of a data item is updated anywhere on the network, the same update is *immediately* applied to all other copies or it is aborted
- Ensures data integrity and minimizes the complexity of knowing where the most recent copy of data is located
- Can result in *slow response time* and *high network usage*
  - the DDBMS spends time checking that an update is accurately and completely propagated across the network.
  - The committed updated record must be identical in all servers

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- Some delay in propagating data updates to remote databases
  - some degree of at least temporary inconsistency is tolerated
  - may be ok it is temporary and well managed
- Acceptable response time
  - updates happen locally and data replicas are synchronized in batches and predetermined intervals
- May be more complex to plan and design
  - need to ensure the right level of data integrity and consistency
- Suits some information systems more than others
  - compare commerce/finance systems with social media

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Advantages and disadvantages of DDBMS
- Distribution, partitioning and replication
- Synchronous vs asynchronous updates
- The CAP theorem

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- NoSQL databases

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder